



Introduction to High Performance Computing

Lecture 02 – Introduction to Message Passing

Holger Fröning
Institut für Technische Informatik
Universität Heidelberg



Concepts of Parallel Architectures



■ Parallel Architectures

- Are pervasive: smartphones, netbooks, notebooks, ...
- Most important resources for HPC: computing units, memory resources, interconnection network
- Most common architecture found today: multi-core systems

■ Multi-core system consists of:

- Multiple computing/processing units, either:
 - Within one processor die (multi-core)
 - Within one computing system (SMP, multiple sockets)
- One single or multiple memory controllers

■ There are any more architecture types:



- Most important: **Shared-Memory** and **Distributed-Memory**



Shared Memory

■ Shared address space

- Each computing unit can access any memory address
- Physical memory may be distributed, access times may vary
 - UMA (uniform memory access)
 - NUMA (non-uniform memory access)
- Any cache can hold copies from any location
 - Cache coherent NUMA (ccNUMA)
 - Usually limited scalability



An address space defines a range of discrete addresses; each address may correspond to a different resource



Shared Memory

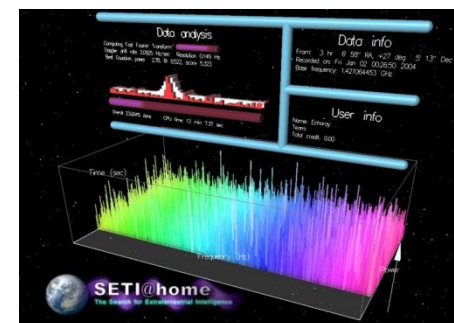

- Shared variables
 - Operations, naming, ordering
- Communication purposes
 - Each unit can read/write any address
 - Used for data
 - Hazards and race conditions require synchronization
- Synchronization purposes
 - Resolve hazards and race conditions
 - Semaphores, condition variables, mutexes, locks
- Implicit communication, explicit synchronization



Distributed Memory

■ Exclusive address spaces

- Operations, naming, ordering
- No (direct) remote access possible (NORMA)
- Physical memory is distributed
- Usually shared memory systems as building blocks
- No global cache coherency required
 - Unlimited scalability (from an architectural point of view)
- Popular example:
SETI@Home





- **Messages for communication purposes**
 - Data transfer (copies)
 - Read example: send a request, receive an answer
 - Write example: send a request, optional answer
- **Messages for synchronization purposes**
 - On the origin node, one process is sending the message
 - On the target node, one process is receiving the message
 - Both will be informed when the other has processed the message
- **Error-prone programming**
- **Programmer is responsible for:**
 - Data distribution
 - Explicit communication
 - Explicit synchronization
 - Many optimizations (huge set of messaging functions, ...)



Shared vs. Distributed Memory

- **Key distinction: number of address spaces**
 - Is (direct) access to the complete memory possible?

| | Shared Memory | Distributed Memory |
|--------------------------|---------------------------|--------------------|
| Number of address spaces | 1 | N |
| Communication | Implicit | Explicit |
| Synchronization | Explicit | Explicit |
| Number of data copies | Typ. 1 | Typ. [1..N] |
| Scalability | Limited (cache coherence) | Unlimited/High |

Synchronization is the most common error source!

Main reason for HPC



Example

- A human discussion as joint work example:
 - Let's assume they use a sheet of paper for the discussion
- If everybody is in the same room, one single sheet of paper can be used
 - Shared resource, but synchronization required
 - Use of pencil and eraser simultaneously, multiple instances...
- If the persons are physically distributed without access to the „local“ sheet, this sheet has to be sent around
 - Messaging
 - Only one copy – weak performance, little synchronization overhead
 - Multiple copies – improved perf., lot's of synchronization
- For both approaches, partitioning is an essential optimization



- HPC favors distributed memory
 - Scalability
- Key tool for HPC today is message passing
- Much more than send and receive
 - Plenty of send variants, plenty of receive variants
 - Collective communication
 - Even communication without saying *receive* (remote memory access, RMA)
 - More in later lectures...



Basics of Message Passing



Parallel Programming

- Nobody wants to write ***N*** programs for ***N*** processes
- Single Program Multiple Data (SPMD) paradigm
 - Write one program, which is internally partitioned to act differently
 - Master/Slave(s), Consumer(s)/Producer(s), Peer model
 - Before execution, each instance is assigned a unique number
 - Let's call this „***rank***“
 - Let's call the total number of ranks „***size***“
 - During execution, the rank determines program behavior
 - Functionality
 - Associated input – better output – data
- SPMD widely used for shared-memory and distributed-memory programming



SPMD Examples

■ Master/Slave

```
if ( rank == 0 ) {
    act_as_master();
} else {
    act_as_slave();
}
```

■ Peer model

```
int work[n];
for ( i = rank/size * n;
      i < (rank+1)/size * n;
      i ++ ) {
    process ( work[i] );
}
```

Note: $(n \% \text{size} == 0)$ is not mandatory



SPMD and Distributed Memory


- Distributed Memory programs run locally with local data
 - Everything explicit
 - Communication
 - Synchronization
 - Data distribution
- } **done using messages**
- Complicated programming
 - Not compatible with shared-memory programming
 - Code porting
 - Development of new algorithms
 - Some help
 - Messaging libraries



Messaging Overview


■ Point-to-point operations

- Data movement between two peers




■ Collective operations

- Collective data movement (among peers or master/slaves)
- Barriers (among peers)



■ Communicators

- Define process groups
- Ranks / size of a communicator
- Every message is assigned to one communicator
- Collective operations refer to all ranks within one communicator





Message Details

■ A message consists of the following

- Source
- Destination
- Tag
- Size
- Data (optional)
- Type

■ Send operation

- Provide all these information

■ Receive operation

1. Receive any message
2. Receive any message from a **given source**
3. Receive any message with a **given tag**
4. Receive any message from a **given source** with a **given tag**



Introduction to MPI



Message Passing Interface (MPI)

- De-facto standard for message passing in HPC
 - MPI 1.0 in 1994, to MPI 2.2 in 2009
- Library for C, Fortran, C++
- <http://mpi-forum.org>
- Plenty of functions – here only a couple
 - Mandatory are about 10-20
- One predefined communicator
 - MPI_COMM_WORLD (containing all processes)



- Always include MPI header file

```
#include <mpi.h>
```

- Initialize library

- Call prior to any other MPI function

```
MPI_Init(&argc, &argv);
```

- Finalize library

- Call after all MPI calls are done

```
MPI_Finalize();
```

- Get rank (unique ID) and size (number of processes) of current process for a given communicator (here default)

```
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
MPI_Comm_size(MPI_COMM_WORLD, &size);
```



„Hello parallel world!“

```
// A very short MPI test routine to show the involved hosts
#include <mpi.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    int size,rank;
    char hostname[50];
    MPI_Init ( &argc, &argv );

    MPI_Comm_rank ( MPI_COMM_WORLD, &rank); // Who am I?
    MPI_Comm_size ( MPI_COMM_WORLD, &size ); // How many processes?

    gethostname (hostname, 50 );
    printf ("Hello Parallel World! I'm process %2d out of %2d (%s)\n",
           rank, size, hostname );

    MPI_Finalize ();
    return 0;
}
```



„Hello parallel world!“

■ Compile:

```
mpicc -Wall mpi_hello.c -o mpi_hello
```

- Any additional compilation parameters (-O3, -L, -l, -h, ...)
- See also `mpic++`, `mpicxx`

■ Prepare execution (only once per user):

- MPI is usually based on ssh-connections. Ensure that no password is required for remote login:

```
ssh-keygen -t dsa -b 1024
```

```
cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys
```



„Hello parallel world!“

- Execute, simplest flavor, 2 processes on one node:

```
mpirun -host creek01[,creek01] -np 2 ./mpi_hello
```

- Execute, 16 processes on two nodes:

```
mpirun -host creek01,creek02 -np 16 ./mpi_hello
```

- Parameters:

- **-np <n>**: start n processes
- **-host**: list of hostnames, used cyclic if more processes than entries in this list
- **-mca ...**: parameters to pass to the MPI library, here use TCP for communication (-mca btl tcp,self)



„Hello parallel world!“

Output:

```
# mpirun -host creek01,creek02 -np 4 ./mpi_hello
Hello Parallel World! I'm process 0 out of 4 (creek01)
Hello Parallel World! I'm process 2 out of 4 (creek01)
Hello Parallel World! I'm process 1 out of 4 (creek02)
Hello Parallel World! I'm process 3 out of 4 (creek02)
```

```
# mpirun -host creek01,creek01,creek02,creek02 -np 4
  ./mpi_hello | sort
Hello Parallel World! I'm process 0 out of 4 (creek01)
Hello Parallel World! I'm process 1 out of 4 (creek01)
Hello Parallel World! I'm process 2 out of 4 (creek02)
Hello Parallel World! I'm process 3 out of 4 (creek02)
```



A first message...

▪ (Blocking) send of a message

```
int MPI_Send(void *buf, int count, MPI_Datatype
datatype, int dest, int tag, MPI_Comm comm)
```

• Example:

```
signal = 42;
```

```
MPI_Send ( &signal, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
```

- A blocking send returns if the provided buffer can be used again.
- Plenty of MPI_Datatypes



A first message...

▪ (Blocking) receive of a message

```
int MPI_Recv(void *buf, int count, MPI_Datatype
datatype, int source, int tag, MPI_Comm comm,
MPI_Status *status)
```

• Example:

```
MPI_Recv ( &signal, 1, MPI_INT, i, 0,
          MPI_COMM_WORLD, &status );
```

▪ A blocking receive returns if a message with the required properties has been received (source, tag)

- MPI_ANY_SOURCE
- MPI_ANY_TAG

▪ Status contains more information about the receive

- MPI_STATUS_IGNORE



Communication example

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char **argv)
{
    int i, size, rank, signal;
    MPI_Status status;
    double starttime, endtime;

    MPI_Init (&argc, &argv);    //Initialisation of MPI
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    if (rank == 0) { //Master: send to every proc and wait for answer
        <see 2nd slide>
    } else { //Clients: wait for message from master and send back
        <see 3rd slide>
    }
    MPI_Finalize(); //Deinitialisation of MPI
    return 0;
}
```



Communication example

```
if (rank == 0) {

    signal = 666;
    starttime=MPI_Wtime();
    for (i=1; i<size; i++) {
        MPI_Send(&signal, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
        printf ("%d: Sent to %d\n", rank, i);
        MPI_Recv (&signal, 1, MPI_INT, i, 0, MPI_COMM_WORLD,
                  &status);
        printf ("%d: Received from %d\n", rank, i);
    }
    endtime=MPI_Wtime();
    printf ("%d: Time passed: %f\n", rank, endtime-starttime);
}
```

```
double MPI_Wtime()
```

Returns time in seconds since an arbitrary time in the past



Communication example

```
//Clients: wait for message from master and send back
else {
    MPI_Recv (&signal, 1, MPI_INT, MPI_ANY_SOURCE, 0,
              MPI_COMM_WORLD, &status);
    printf ("%d: Received from %d\n", rank, status.MPI_SOURCE);
    MPI_Send(&signal, 1, MPI_INT, status.MPI_SOURCE, 0,
             MPI_COMM_WORLD);
    printf ("%d: Sent to %d\n", rank, status.MPI_SOURCE);
}
```



Blocking vs. Non-blocking

■ Blocking

- Returns after send or receive complete
- What means complete?

■ Non-blocking variants

- Guarantees nothing, may return immediately
- Used for improved overlap between computation and communication

```
int MPI_Isend(void *buf, int count, MPI_Datatype  
datatype, int dest, int tag, MPI_Comm comm,  
MPI_Request *request)
```

```
int MPI_Irecv(void *buf, int count, MPI_Datatype  
datatype, int source, int tag, MPI_Comm comm,  
MPI_Request *request)
```



Non-blocking: test/wait for completion

■ Blocking: Wait

```
int MPI_Wait(MPI_Request *request, MPI_Status  
*status)
```

- Guarantees that request has been processed

■ Non-Blocking: Test

```
int MPI_Test(MPI_Request *request, int *flag,  
MPI_Status *status)
```

- Updates the request, but does not ensure completion



Questions?

Next lecture:

Basics