



Introduction to High Performance Computing

*Lecture 03 – Applications, Performance Increase,
Top500*

Holger Fröning
Institut für Technische Informatik
Universität Heidelberg




Example Applications Fields




Example application fields


- **Oil and gas**
 - Seismic processing
- **Financial services**
 - Automated option pricing and trading
- **Bioscience**
 - Genetic sequencing and chemistry
- **Government**
 - Searching and encryption engines
- **Digital content creation**
 - Movie animation
- **Scientific research**
 - Astrophysics, particle physics
 - Biology: molecular dynamics
- **Industry**
 - Fluid dynamics
- **Meteorology & Climatology**
- **Electronic design automation**
 - Chip design
- **Finite element methods**
 - Crash simulations
- ...




CFD



Molecular Dynamics (MD)



Cosmology



Weather/Climate Research




A more abstract view...

■ Basic operation types in HPC

- Complex processing of regular data structures
 - Vectors
 - Arrays
 - Elements
- High degree of parallelism

■ Trivial example: Matrix addition

- “Domain decomposition”






Simple parallelization example


Serial execution

- No communication
- Sequential processing of elements



Parallel execution

- Communication?
- Synchronization?
- Parallel processing of (blocks of) elements





Molecular dynamics


■ Motivation

- Protein Folding
- Digital simulation instead of biochemics
- Computationally intensive
- Runtime of several months
- STMV: 160 genes, 100ns/day for Petascale-class


■ Goal: Calculation of the molecule's shape

- Double precision floating point
- Calculation of forces in between the atoms of the molecule and the surrounding
 - Forces: Electrostatic (Coulomb) & Van der Waal
- Time step = 1 femto (10^{-15}) second


■ “N-body problem”



nasa.gov



wikipedia.org






Molecular dynamics: N-body

- N-body problem
 - For each atom in a 3D system

```

Do repeat
  Increase time step t
  Foreach atom i
    Foreach atom j (j != i)
      Compute force(s) from j to i
    Sum all forces on i
  Next j
  Compute next position of i
Next i
Repeat until stable
    
```




forces $\sim 1/d^2$
 $\sim 82\%$ time
 ~ 52 FP ops

- No special treatment of borders here...
 - Approx. 60 variants



- Genetics related research
- String distance computation
- Using Smith-Waterman
 - Exact string matching algorithm
 - Finds optimal local alignment
 - Computes a matching score $H(i,j)$ of two input strings \mathbf{S} and \mathbf{T} using a 2D matrix
- Other applications:
 - Motif discovery, data mining




$$H(i,j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right. \begin{array}{l} \text{Match/Mismatch} \\ \text{Deletion} \\ \text{Insertion} \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n.$$




Computational Fluid Dynamics (CFD)

- Numerical fluid simulations
- Lattice-Boltzmann method (LBM)
 - Simulation at particle scope
 - Discretization by grid, at microscopic level solution of mutual reaction as described by the Boltzmann equation
- Boltzmann equation in the case of large average free path lengths
 - Currents in (depleted) gases
 - Neutron distribution in atomic reactors
- Otherwise: E.g. Navier-Stokes equation
 - Much simpler
 - Current in liquids
- Rather few communication, large messages between pairs of two



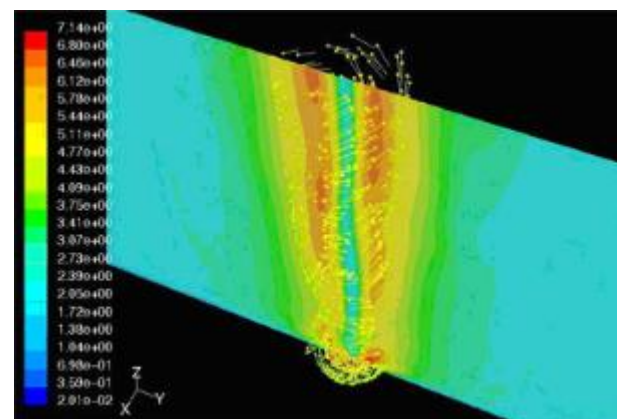
ucar.edu



■ See CFD

■ Examples


- Hurricane prediction
- Tornado simulation
- Local weather forecast
- Global Climate Modeling






Implementations

- Solving/Approximation of Partial Differential Equations (PDE) or CFD
- Stencil codes
 - Iterative kernels
 - Regular, invariable structure
- Finite Element Methods (FEM)
 - Irregular structures
 - For complex or variant structures
 - Different accuracies
 - Adaptive Mesh Refinement (AMR)
- Technical and scientific computing is mostly modeling
 - Based on time steps
 - Iterative solution until results are stable or simulation period is over



6-point 3D stencil, courtesy: wikipedia.org





Basics of Performance Increase

*Performance Increase of
Technologies and Applications*

Moore vs. Amdahl



Moore's Law


■ Gordon Moore

- 1965: Doubling each year
- 1975: Transistor count of ICs doubling every two years

■ Derived "laws"

- CPU performance doubling every 18 months
- Memory size four times every three years
- Memory performance doubling every 10 years
- At the same costs double performance every two years

Relative
Cost/Component



Gordon Moore,
Mitbegründer Intel


intel.com

Number of
components per IC



Moore's Law


wikipedia.org





Moore's Law

- Industry is trying to keep the pace
 - Self-fulfilling prophecy
 - “positive feedback between belief and behavior”
- Atoms as fundamental lower bound
 - Even then, increase of die size can maintain the law
 - Intel's statements about end of Moore's law
 - 2003: 2013-2018
 - 2005: until 2015
 - 2008: until 2029
- Bernie Meyerson (IBM):
7-9nm is the limit
 - Quantum mechanics effects





- Speed-up: „How much faster can one program be executed“
- Assumption: instead of one resource, N identical resources are available
- Naive: More resources, faster execution
- A bit more realistic: N resources yield an execution time of $1/N$
 - No overhead assumed
- Reality: significant loss
 - Break-even point when execution time starts to increase again



Speed-up - Definitions

■ For a given algorithm:

- $\text{SerTime}(n)$ = time of the best serial implementation for an input of size n
- $\text{ParTime}(n,p)$ = time of the parallel implementation, using p parallel computing units

■ Sanity check: $\text{SerTime}(n) \geq \text{ParTime}(n,1)$

- The other case is not uncommon

■ Speed-up:

- $\text{Speedup}(p) = \text{SerTime}(n) / \text{ParTime}(n,p)$
- $\text{Efficiency}(p) = \text{SerTime}(n) / (p * \text{ParTime}(n,p))$




Speed-up - Notes

- $1 \leq \text{Speedup}(p) \leq p$
- $0 \leq \text{Efficiency}(p) \leq 1$

- Linear speed-up:
 $\text{Speedup}(p) = p$


- Superlinear speed-up:
 $\text{Speedup}(p) > p$
 - Usually not possible






Speed-up – Real experiment

Speed-up of HPCC RandomAccess



Speed-up of HPCC RandomAccess





Amdahl's Law

- Model to find the maximum improvement in terms of performance
 - Assumption: only a fraction of the runtime can be parallelized (parallel fraction **P**).
 - Assumption that the other fraction is the serial one: serial fraction **S**
 - Then: **P** + **S** = 1
 - As fraction **P** is processed in parallel, this fraction of time is reduced (**N** parallel execution units)

$$Speedup = \frac{1}{(1 - P) + \frac{P}{N}}$$


- Notes:
 - Speed-up has an upper limit dependent on **S**, not on **N**!



Amdahl's Law

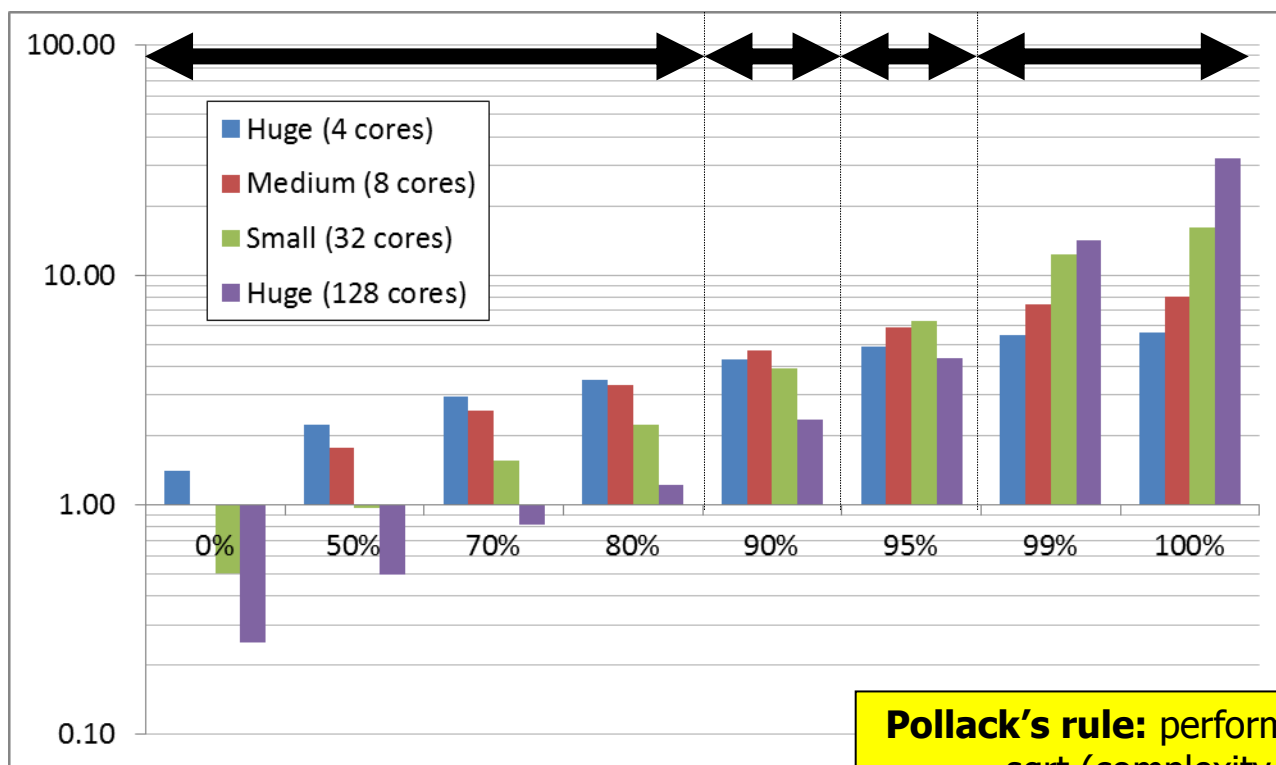
$$Speedup = \frac{1}{(1-P) + \frac{P}{N}}$$

Amdahl's Law
(dependant on parallel portion P)





Amdahl's Law - Implications



Assumptions
80W &
200mm² for
cores

Pollack's rule: performance increase \sim
sqrt (complexity increase)

Number of Cores	4	8	32	128
Power (W) / Core	20	10	2,5	0,6
Area (mm ²) / Core	50	25	6	1,5
Relative Performance R	140 %	100 %	50 %	25 %



■ Amdahl himself ...

1. ... wanted to claim that **parallel computing is not viable**
 - "Validity of the Single Processor Approach to Achieving Large-Scale Computing Capabilities", *AFIPS Conference Proceedings*, 1967.
2. ... was an **optimist**
 - Extra work is required for parallelization
 - Synchronization, communication, management, ...
 - In this regard his law is too optimistic
3. ... was a **pessimist**
 - We can (have to?) scale the problem size with **N**
 - Gustafson's law – superlinear speedup (1988)
 - Parallel algorithms exist that reduce fraction **S**
 - Superlinear speed-up due to caching effects



Performance increase - Summary

- Increase of performance according to Moore's Law
 - Technology 😊
- Increase of performance according to Amdahl's Law
 - Limited by serial fraction ☹️

"Everyone knows Amdahl's Law, but quickly forgets"
Dr. Tom Puzak, IBM Research, 2007

- Sources for serial fraction
 - Data dependencies
 - Communication & synchronisation is costly
- ⇒ Optimize these components ☹️
- ⇒ Increase problem size 😊
- Increases percental fraction of ***P***



TOP500 List



- <http://www.top500.org>
- Biannual list: 500 fastest computer systems world-wide
 - LinPACK benchmark
 - „Dense system of linear equations”
 - $\frac{2}{3} n^3 + O(n^2)$ double precision floating point operations
 - Highly scalable, problem size can be chosen arbitrary
- Computationally intensive, not memory-bound
 - Little requirements on memory bandwidth and capacity
- Old lists available on-line
 - History and trends




TOP500 – List of 06/2014

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	National Super Computer Center in Guangzhou China	Tianhe-2 (MilkyWay-2) - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P NUDT	3,120,000	33,862.7	54,902.4	17,808
2	DOE/SC/Oak Ridge National Laboratory United States	Titan - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x Cray Inc.	560,640	17,590.0	27,112.5	8,209
3	DOE/NNSA/LLNL United States	Sequoia - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom IBM	1,572,864	17,173.2	20,132.7	7,890
4	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect Fujitsu	705,024	10,510.0	11,280.4	12,660
5	DOE/SC/Argonne National Laboratory United States	Mira - BlueGene/Q, Power BQC 16C 1.60GHz, Custom IBM	786,432	8,586.6	10,066.3	3,945
6	Swiss National Supercomputing Centre (CSCS) Switzerland	Piz Daint - Cray XC30, Xeon E5-2670 8C 2.600GHz, Aries interconnect , NVIDIA K20x Cray Inc.	115,984	6,271.0	7,788.9	2,325
7	Texas Advanced Computing Center/Univ. of Texas United States	Stampede - PowerEdge C8220, Xeon E5-2680 8C 2.700GHz, Infiniband FDR, Intel Xeon Phi SE10P Dell	462,462	5,168.1	8,520.1	4,510
8	Forschungszentrum Juelich (FZJ) Germany	JUQUEEN - BlueGene/Q, Power BQC 16C 1.600GHz, Custom Interconnect IBM	458,752	5,008.9	5,872.0	2,301
9	DOE/NNSA/LLNL United States	Vulcan - BlueGene/Q, Power BQC 16C 1.600GHz, Custom Interconnect IBM	393,216	4,293.3	5,033.2	1,972
10	Government United States	Cray XC30, Intel Xeon E5-2697v2 12C 2.7GHz, Aries interconnect Cray Inc.	225,984	3,143.5	4,881.3	



TOP500 – Architecture Share

Architecture Share Over Time
1993-2011

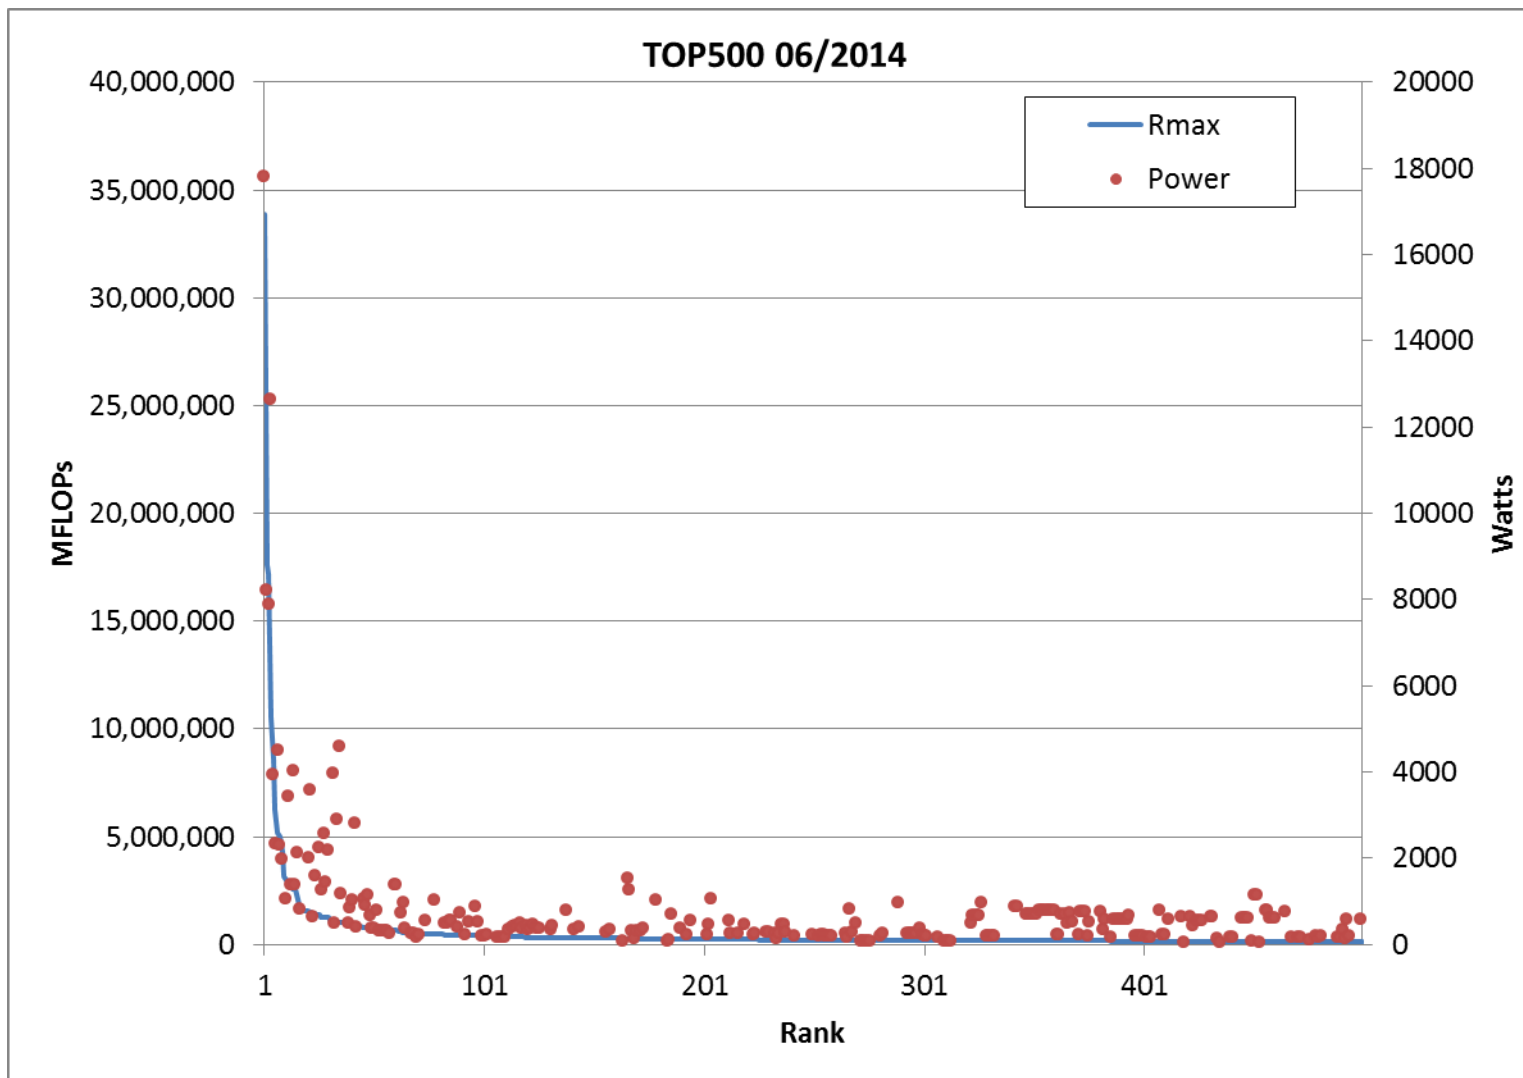


N: number of nodes
C: number of cores per node

- System with N nodes, each C cores
 - 1 node equals 1 address space
- Massively Parallel Processors (MPP)
 - $N > C$, $N \gg 1$
(whatever that means)
- Cluster
 - $N > C$, $N > 1$
- Symmetric Multi-Processors (SMP)
 - $N = 1$
- Constellations
 - $N < C$, $N > 1$
- Single Processor
 - $N = C = 1$
- Others
 - Never seen




TOP500 – Performance & Power



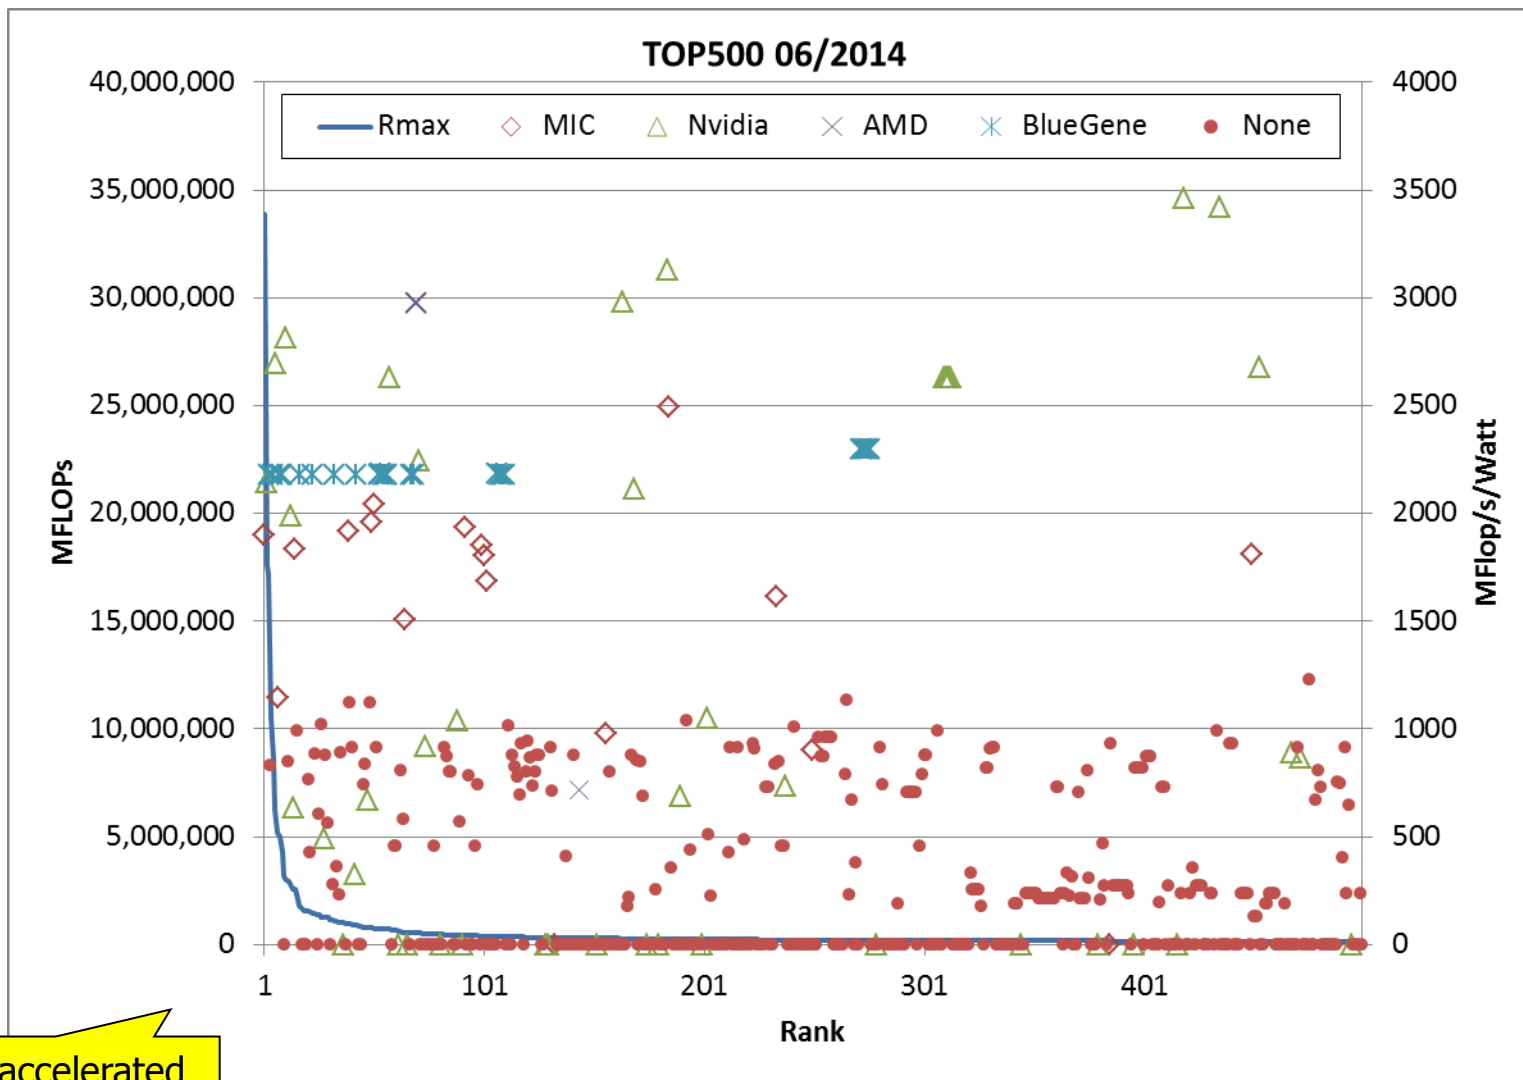


TOP500 – Power Efficiency





TOP500 – Accelerated Systems





TOP500 - Summary

- Excellent tool for trend analysis
 - Introductions, motivating data
- Maybe too limited due to the single workload
 - No scalability worries
- Alternatives
 - Graph500:
<http://www.graph500.org>
 - Green500:
<http://www.green500.org>
 - HPC-Challenge:
<http://icl.cs.utk.edu/hpcc>

- Exascale at 2GFLOPs/Watt (BlueGene):
 - 1,000,000,000,000,000,000 Flops
(1,000,000,000 GFlops)
 - 500,000,000 Watts
(500 MWatt)
 - @50MWatt: 20GFLOPs/Watt required
- ➔ Extreme specialization
 - What about too specialized?