



Introduction to High Performance Computing

Lecture 04 – Parallel Computing

Holger Fröning
Institut für Technische Informatik
Universität Heidelberg



Parallelism



- Sequential vs. parallel processing completely different
- Multi-/Many-core era
 - Applications designed for single-core
 - **Concurrency** is fundamental for algorithms and applications
- Number of cores/CPU increasing
 - **Scalability** also fundamental
- Further motivations:
 - Performance increase, distributed systems, tolerating I/O Blocking


Parallel programming: Concurrency & Scalability (I & II)



Concurrency

Sequential Program

- Single thread of control
- Instructions executed sequentially



Concurrent Program

- Several autonomous sequential threads
- Parallel execution
- Execution determined by implementation


Implementations

- Multi-programming
 - Executing multiple threads on one single resource
 - Time-multiplexing
- Multi-processing
 - Executing multiple threads on one multiple resources
 - Multi-processor, Multi-core
- Distributed processing
 - Executing multiple threads on one multiple independent resources
 - Cluster, Grid, Cloud



Concurrency vs. Parallelism

- **Concurrency is not (only) parallelism!**
- **Concurrency by interleaving**
 - Only logical “parallel” execution on one single resource
 - Appearance of “simultaneous” execution
- **Parallelism**
 - True parallel, simultaneous execution
 - Requires several, parallel resources
- **Example for concurrency:**
 - Multiple ATMs (“EC-Automaten”) and account balance



Error-free execution on sequential hardware not necessarily implies error-free execution on parallel hardware



Levels of Parallelism – Traditional Approach

▪ Program level

- Coarse grained
- Concurrent execution of multiple programs, or of a single program with different input data sets

▪ Procedure level

- Medium grained
- Different parts of a program are executed concurrently on different parts of a computing system, or one single part with different input data sets

▪ Instruction level

- Fine grained
- Concurrent computation of multiple variables in one procedure

▪ Microcode level

- More fine grained
- Instruction → Operations (Phases)
- Execution of different phases of different instruction in different pipeline stages or superscalar execution units simultaneously

▪ Bit level

- Extreme fine grained
- Processing of words only, consisting of multiple bits



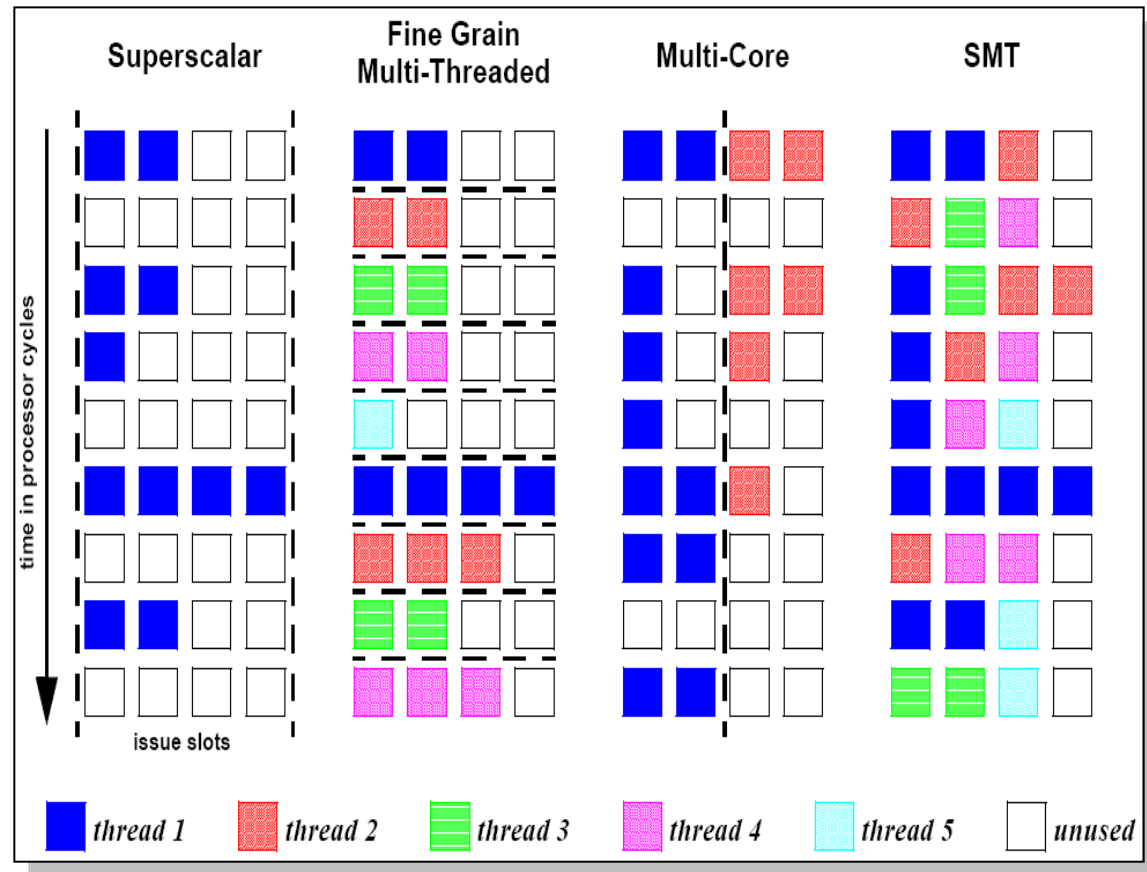
Levels of Parallelism – Modern Approach

- **Instruction Level Parallelism (ILP)**
 - Parallelism of one instruction stream
 - Huge amount of dependencies and branches
 - Limited parallelism (~4-6)
- **Thread Level Parallelism (TLP)**
 - Parallelism of multiple independent instruction streams
 - Less amount of dependencies, no limitations due to branches
 - Limited by maximal concurrently executable I-streams
- **Data Level Parallelism (DLP)**
 - Vectorization techniques
 - Applying one operation on multiple elements of a data structure
 - Parallelism dependent on data structure



Levels of Parallelism (3)

- Exploiting parallelism in different CPU architectures
 - ILP
 - TLP
 - Why no DLP?
- What about GPUs?





Computing Model



Computing model (1)

■ von-Neumann architecture

• Main units

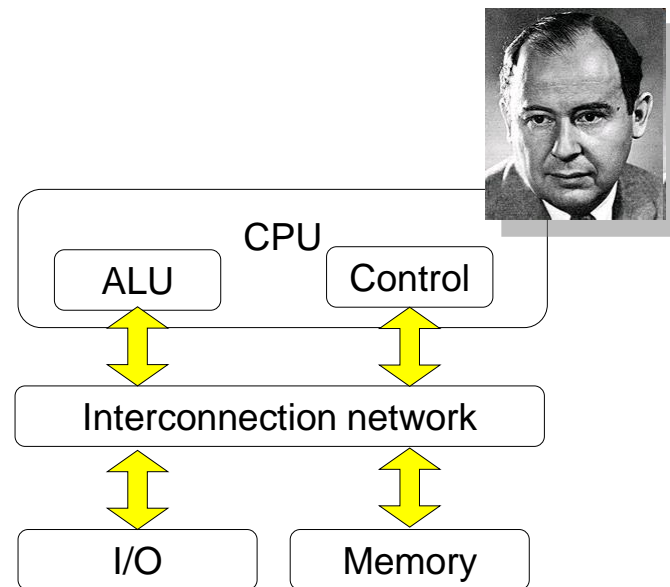
- CPU (Control & Compute)
- I/O
- Memory

• “Node” for HPC systems

■ von-Neumann bottleneck

- ALU faster than memory
- Costs for control and communication higher than computing costs

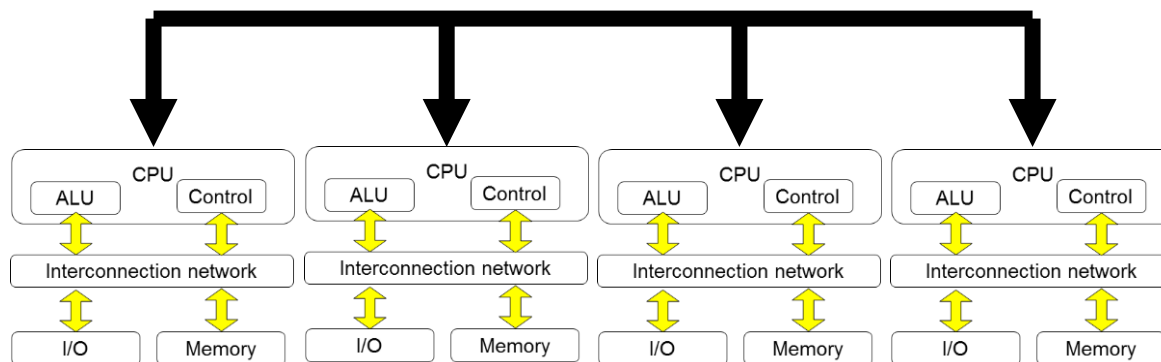
■ Harvard Architecture: Separation of data and instruction memory





Computing model (2)

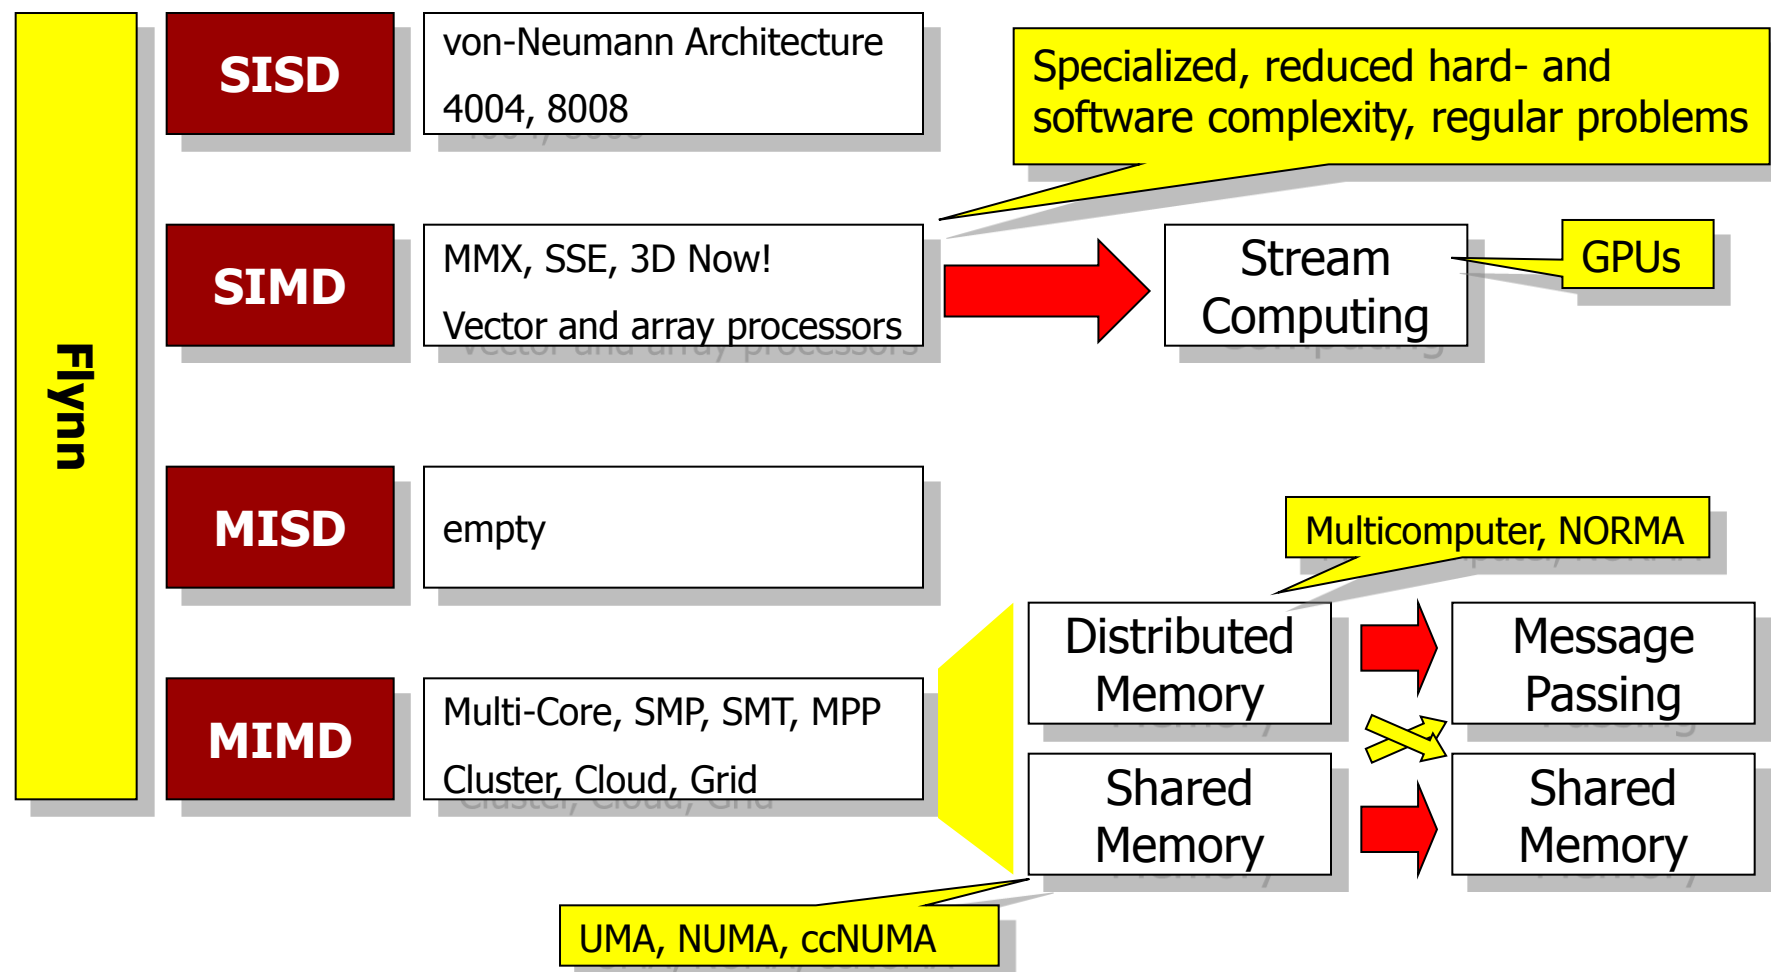
- Multicomputer
 - Multiple nodes
 - Interconnection network
- Many parallel instruction streams
- Local (and remote) memory accesses
- According to Flynn?



Parallel programming: Locality (III)



Classification by Flynn






Distributed and Shared Memory


■ Shared Memory

- Shared use of one copy
- Scalability issues
- Atomicity, locking, synchronization



■ Distributed Memory

- Explicit data exchange
- Only access to local memory
- Data distribution and communication scheme





1. Concurrency

- Functional Decomposition, Domain Decomposition, Pipeline Decomposition
- Re-engineering for parallelism
 - Control dependencies, data dependencies

2. Parallel programming paradigms

- Shared memory: PThreads, OpenMP
- Distributed memory: Message-passing
- Data parallel operations (SIMT): CUDA, OpenCL

3. Supporting structures

- SPMD, loops, master/worker, fork/join, data structures



- von-Neumann:
 - 1 node → 1 instruction stream
- Multicomputer:
 - n nodes → n instruction streams
 - Too complex
- Modular approach
 - Simple components made of abstract elements
 - Data structures, loops, procedures
- Single Program Multiple Data (SPMD)

Parallel Programming: Modularity (IV)



■ Used for this lecture: tasks & Channels

• Task


- Computations
- Instructions & Memory

• Channel

- Communication among tasks
- Message-based
- Blocking receives

■ Computation & Communication

■ Data dependencies





■ Message Passing

- Difference:
 - Message Passing: send to x
 - Task/Channel: send over channel y
- SPMD

■ Data Parallelism

- Applying one operation to multiple elements of a data structure
- SPMD

■ Shared Memory


- Uniform memory access from user's point of view
 - No explicit communication
 - Locks & Semaphores
- SPMD



Synchronization

Synchronization is the enforcement of a defined logical order between events. This establishes a defined time-relation between distinct places, thus defining their behavior in time.

- Communication & synchronization
 - Explicit / implicit
- SIMD: one instruction stream, no synchronization necessary
- MIMD: synchronisation necessary
 - Shared variables
 - Process synchronization
 - Blocking message exchange







Algorithm Design



Book is online at:
<http://www.mcs.anl.gov/~itf/dbpp>

■ Foster's PCAM

- Partition
- Communicate
- Agglomerate
- Map





Number of Tasks \gg Number of P

■ PCAM: Partitioning




- Ignore technical aspects like number of processing units
- Maximal granularity
- Partition computation and data
 - Domain Decomposition
 - Functional Decomposition
 - Pipeline Decomposition
- Avoid replication, disjoint partitioning
 - See also minimization of communication

➔ Partitioning



▪ PCAM: Partitioning

- Example 1: Domain Decomposition
- Typical uses: data parallelism, e.g. arrays & trees






▪ PCAM: Partitioning

- Example 2: Functional Decomposition
- Typical uses:
 - Function calls
 - Different loop iterations
- Rather too many tasks than too few!


Climate Computing Model





▪ PCAM: Partitioning

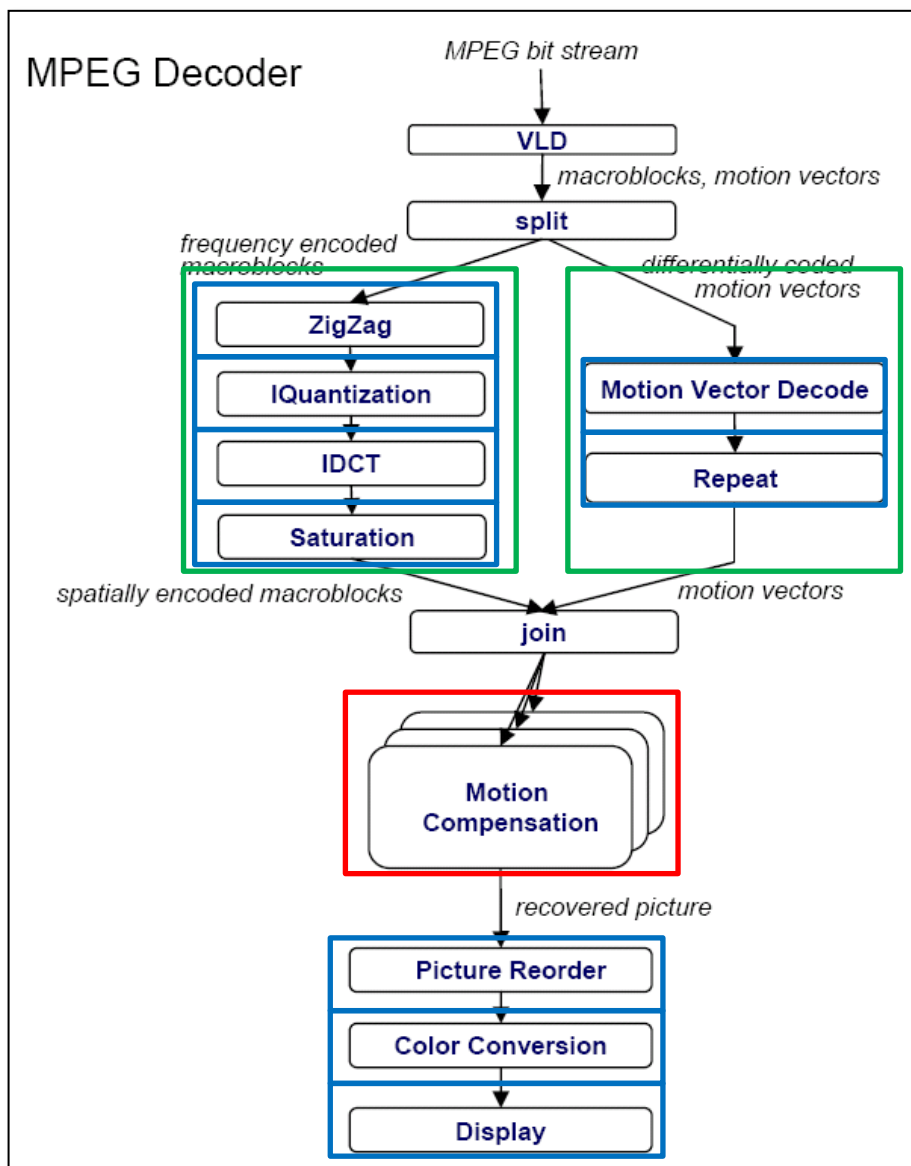
- Example 3: Pipeline Decomposition
- Data flow through several pipeline stages
- Typical uses:
 - Instruction pipelining in modern CPUs





Algorithm Design

- Identify possible decomposition techniques!
- Domain Decomposition
 - Red
- Functional Decomposition
 - Green
- Pipeline Decomposition
 - Blue





- **PCAM: Communicate**
 - Execution of partitions **concurrently**, but not **independently**
 - Data dependencies → communication & synchronization
- Complex for DD, rather simple for FD
- Local/global, structured/unstructured, static/dynamic, synchronous/asynchronous

→ Communication scheme

- **Data-parallel language**
 - Requires data-parallel operations and data distribution. Channels actually not necessary, but help for locality and communication costs



■ PCAM: Communicate

- Example for local communication: stencil operation
 - Simple numerical computation: finite difference method (iterative method used to solve a linear system of equations)
 - **Gauss-Seidel (GS)**

$$X_{i,j}^{(t+1)} = \frac{4X_{i,j}^{(t)} + X_{i-1,j}^{(t+1)} + X_{i+1,j}^{(t)} + X_{i,j-1}^{(t+1)} + X_{i,j+1}^{(t)}}{8}$$

vs. **Jacobi**

$$X_{i,j}^{(t+1)} = \frac{4X_{i,j}^{(t)} + X_{i-1,j}^{(t)} + X_{i+1,j}^{(t)} + X_{i,j-1}^{(t)} + X_{i,j+1}^{(t)}}{8}$$

- GS optimal for sequential execution (fewer iterations), but too many dependencies for parallel execution
- GS execution: diagonal wave front or Red/Black method



▪ PCAM: Communicate

- Global communication
 - E.g. global addition (parallel reduction)

$$S = \sum_{i=0}^{N-1} X_i$$

- Cons: **O(N)**, **centralized & sequential**
- More equal distribution of computation and communication, **O(N-1)**

$$S_i = X_i + S_{i-1}$$

▪ Divide & Conquer to exploit parallelism

- Tree structures, as long as partitions can be computed independently
- Associativity of addition, **O(log N)**



■ PCAM: Agglomeration

- From the abstract to the concrete
- Fixing the parallel computer model

■ Goal

- Increase granularity (coarse-grain)
- Maintaining flexibility, therefore reducing development costs


■ Number of tasks $T \geq$ number of processors P

- Depending on use case:
 - One order of magnitude more T s than P s (parallel slackness)
 - $T == P$ (HPC)
- SIMD: $T = 1$
- If $T = P$, then mapping (almost) done



■ PCAM: Agglomeration

- Combining of tasks
 - Increase of granularity
- Motivation: Reducing communication costs
 - Fixed & variable fraction (surface-to-volume effects)





▪ PCAM: Agglomeration

- Replication of data and computation

▪ Example: global sum

- Chained: $2(N-1)$ steps (sum & broadcast)
 - ➔ Redundant computation in a ring, no broadcast ($(N-1)$)
- Tree-based: $2 \log N$ steps (sum & broadcast)
 - ➔ Redundant computation in a butterfly, no broadcast ($\log N$)

➔ Reducing Communication



■ **PCAM**: Mapping

- Assignment: task \leftrightarrow processor & memory
 - Place tasks that can execute concurrently on different processors
 - Place tasks that communicate frequently on the same processor
 - Note that this implies conflicts
- Mapping not necessary for:
 - Uni-processors or shared memory systems with automatic mapping
 - Hardware mechanism or the OS responsible for scheduling



■ Mapping problem is NP-complete

■ Dynamic Load Balancing



■ PCAM: Mapping

1. Concurrent tasks on different Ps
2. Frequently communicating tasks on same P






Ready to run!



Parallel Computing

Algorithm Design - Examples

- Matrix multiply
- Stencil operation
- N-Body problem






Parallel Computing

Algorithm Design - Examples

■ Stencil codes (e.g. Jacobi method)



- Approximation by time steps



$$X_{i,j}^{(t+1)} = \frac{4X_{i,j}^{(t)} + X_{i-1,j}^{(t)} + X_{i+1,j}^{(t)} + X_{i,j-1}^{(t)} + X_{i,j+1}^{(t)}}{8}$$

■ N-Body codes

- Gravitational forces, electrostatical forces
- Smoothed particle hydrodynamics (simulating fluid flows)
- Superposition
- Approximation by time steps





Parallel Computing Summary

- **Concurrency and parallelism of fundamental importance**
 - Granularity
 - ILP, TLP, DLP
- **Characteristics of “good” parallel programs**
 - Concurrency, Scalability, Locality and Modularity
- **Algorithm design**
 - Partition, Communicate, Agglomerate, Map
- **Parallel computing highly dependent on architecture!**
 - (Flynn’s classification,) shared & distributed memory
- **Literature**
 - Foster Online
 - <http://www.mcs.anl.gov/~itf/dbpp>
 - Introduction to Parallel Computing
 - <http://www-users.cs.umn.edu/~karypis/parbook>