



Introduction to High Performance Computing

Lecture 05 – Practical Parallel Programming Example

Holger Fröning
Institut für Technische Informatik
Universität Heidelberg



Matrix Multiply

✦ Why always Matrix Multiply?

- One of the most heavily optimized codes in HPC
- Interesting access patterns
- Good mixture of sufficient complexity but still simple enough for a comprehensive understanding
- Finally, it's an important operation!

✦ Used in many applications as computational kernel


- In particular for sparse matrix operations

✦ Here: for dense matrices

- Experiments and learning
- High sustained/peak ratio
- Test system/compiler/OS

✦ Note on notation

- $M[\text{row}, \text{column}] = M[\text{row}][\text{column}]$
- Analogous to C





Matrix Multiply – Sequential version

✦ $C = A[][] * B[][]$


- $A[i,j] == A[i][j]$

✦ Locality

- Row-major order storage for C programming language

✦ Caching effects

- Nice for A, bad for B
- True/False sharing
- Potentially evicts other useful blocks (3C)
 - Compulsory
 - Conflict
 - Capacity



```
for i in n
  for j in n
    for k in n
      C[i,j] = C[i,j] + A[i,k]*B[k,j]
```



Matrix Multiply - Analysis

✦ Analysis

- Assume square matrices
- Assume perfect write-through cache (no conflict, no capacity)

✦ Number of flops: $f = 2 \cdot N^3$

- N^2 elements in C, each N steps
- Each step: multiply & add

✦ Number of cache misses:


$$m = 3 \cdot N^2$$

- Load from A, B, C, store to C
- Counting all accesses: $4 \cdot N^3$

✦ Ratio mem/flop:

$$r = m/f = 2/N = O(1/N)$$

- Computationally intensive
- Peak performance expected for cache-based processor

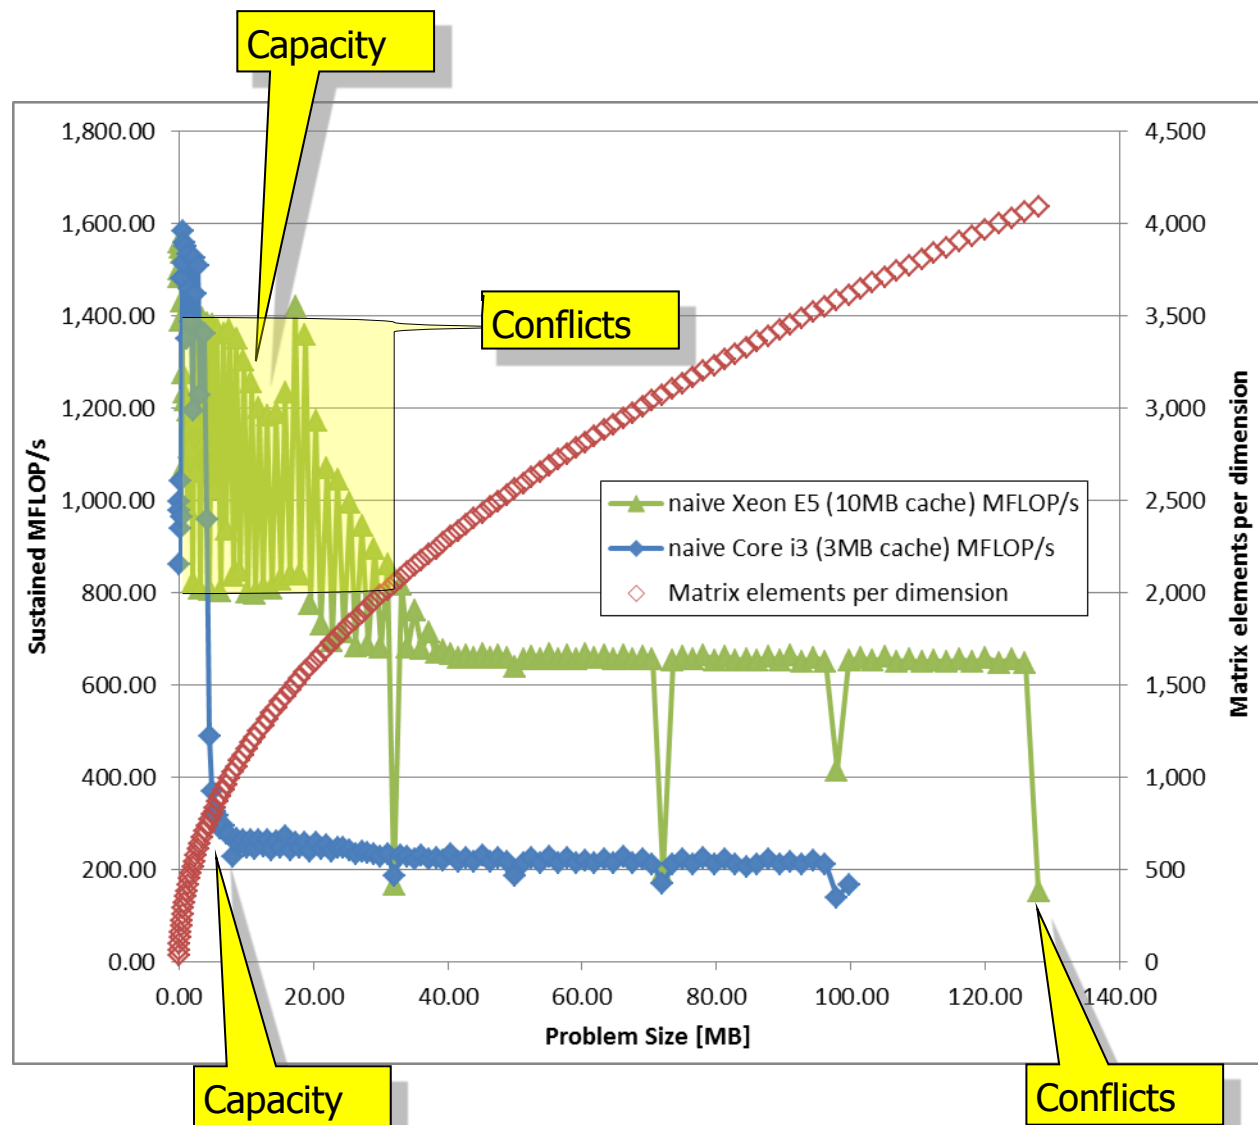


```
for i in n
  for j in n
    for k in n
      C[i,j] = C[i,j] + A[i,k]*B[k,j]
```



Matrix Multiply - Baseline Performance

- ✦ Core i3, 3.1GHz
dual core, AVX
 - 49.6 GFLOP/s peak
- ✦ Xeon E5, 2.4GHz
quad core, AVX
 - 76.8 GFLOP/s peak
- ✦ Size = 3 matrices
each N^2
- ✦ Flops = $2N^3$
- ✦ Nice caching
effects visible
 - Core i3
 - Xeon E5





Matrix Multiply – Analysis


✦ Suffers from non-linear access to B

- „Pseudo-Random“ from the cache point of view
- Neither spatial nor temporal locality exploitable

✦ $r = O(1/N)$ only holds for optimal caching

✦ Otherwise:

- $m_{\text{uncached}} = 3 \cdot N^3$ for all accesses (worst caching)
- $F = 2 \cdot N^3$
- $r = m/f = 2$ or $O(1)$



```
for i in n
  for j in n
    for k in n
      C[i,j] = C[i,j] + A[i,k]*B[k,j]
```



Matrix Multiply – Seq. Version First Optimization

✦Solution: transpose B




- In a linear load fashion, written back with stride
- Optimized for caches

✦Costs:

- $f_t = 0$
- $m_t = 2N^2$

✦New overall r:

- $r = (m_{\text{cached}} + m_t) / (f + f_t)$
 $= (3N^2 + 2N^2) / (2 \cdot N^3 + 0)$
 $= 5/(2N)$
 $= O(1/N)$






Matrix Multiply – Seq. Version First Optimization

✦ Theoretical peak in this case (single threaded, no AVX or SSE)

- Core i3: 3.1×2 (MADD) = 6.2 GFLOP/s





Tiling/Blocking

✦ Increase locality by reordering memory accesses

- Associativity
- Tiling or blocking
- Each TxT tile uses each element T times

✦ Calculate only parts of the elements of C, so that access pattern has high locality


- Beneficial for both sequential and parallel algorithms

Time ↓

$C_{0,0}$	$C_{1,0}$	$C_{0,1}$	$C_{1,1}$
$B_{0,0} * A_{0,0}$	$B_{0,0} * A_{1,0}$	$B_{0,1} * A_{0,0}$	$B_{0,1} * A_{1,0}$
$B_{1,0} * A_{0,1}$	$B_{1,0} * A_{1,1}$	$B_{1,1} * A_{0,1}$	$B_{1,1} * A_{1,1}$
$B_{2,0} * A_{0,2}$	$B_{2,0} * A_{1,2}$	$B_{2,1} * A_{0,2}$	$B_{2,1} * A_{1,2}$
$B_{3,0} * A_{0,3}$	$B_{3,0} * A_{1,3}$	$B_{3,1} * A_{0,3}$	$B_{3,1} * A_{1,3}$



Commercial BLAS library



Source: <http://software.intel.com/en-us/articles/intel-mkl>
Per core: achieved about 25 GFLOP/s out of 27.2 GFLOP/s peak



Parallelization



✦ **Now:** local and remote accesses

- Distinguish between local accesses m_l and remote accesses m_r
- Remote access much worse than local access

✦ **Serial version:**

- Due to memory gap, f/m crucial for overall performance

✦ **Parallel version:**

- Hierarchical approach, first optimize m_r , then m_l
- Remote communication more important for overall performance



Parallel MMULT - Parallelization

✦PCAM: Partition, Communicate, Agglomerate, Map

✦Partition

- Domain decomposition of C (output matrix)
- Each $C[i,j]$ assigned to one task

✦Resulting **Communication** pattern depends on distribution of input matrices A,B

1. Copies: no communication except for initialization/completion
 - For large problem sizes not possible (capacity constraints)
 - Huge initialization overhead, $2N^2$ elements to copy
2. Partitioning of A,B according to C: same as unique memory references ($m = 4N^2$), but here m = number of (potential) messages
 - Depends on tiling/blocking

✦Agglomeration

- Increase computational load per thread/process
- Reduce communication overhead

✦Mapping

- Easy as communication pattern is rather simple (global communication dominates)



✦ Overlap between computation and communication

- For complete matrix copies, dependencies will prevent overlap
 - 3 clearly separated steps: distribute, compute, collect
- For partitioning, fine grain communication allows for overlap:
 - While the first blocks are processed, further distribution of input matrices can already take place

✦ Goal is provide as much overlap as possible

- Latency hiding resp. hiding of communication costs
- Blocking/Tiling



✦ Data distribution primary concern for agglomeration

- 1D partitioning by blocks of rows or columns
- 2D partitioning by both blocks of rows and columns

✦ Optimal choice depends on algorithm

- Locality, communication pattern and dependencies


✦ Here: 1D (row and column) and 2D analyzed

✦ In general for MPI: owner computes and handles associated communication



Parallel MMULT – 1D row partitioning

- ✦ Data layout: process i owns row(s) $C[i, *]$ (short: $C[i,]$)
- ✦ Owner computes: process i calculates $C[i,]$
- ✦ $C[i,] = A[i,] * B[,]$
 - B has to be owned completely, but only parts of A and C

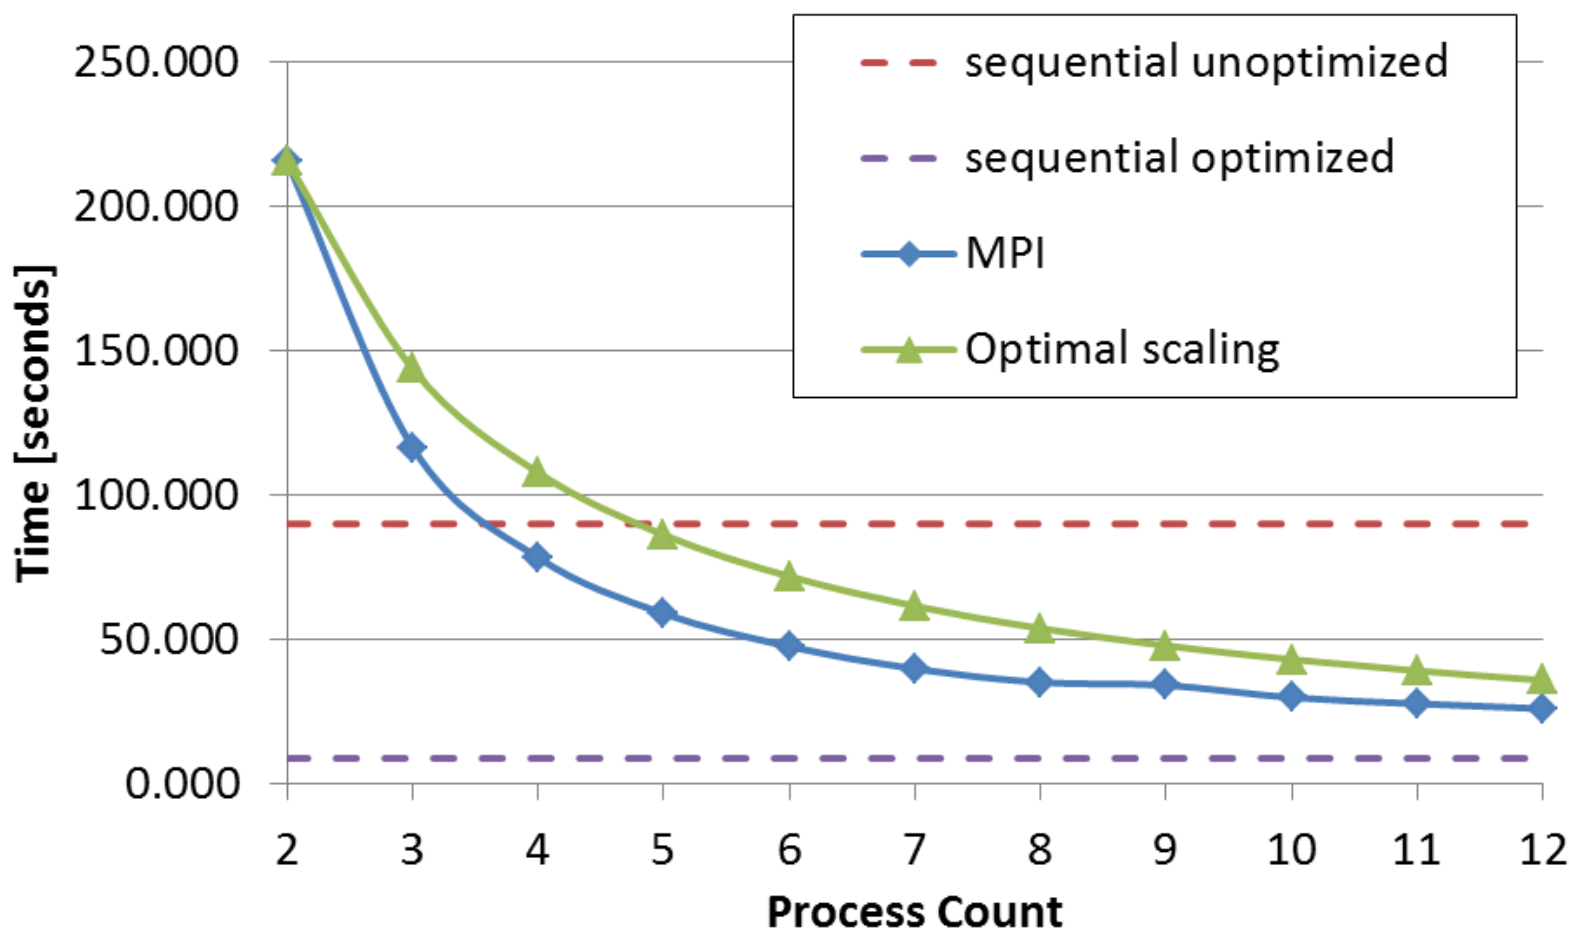




Parallel MMULT – 1D row partitioning

Older experiment!

MPI Matrix Multiply on 6 Dual-Core Nodes






Parallel MMULT – 1D row partitioning

✦ Analysis

- + No communication among workers
- + Nice scalability
- Large memory footprint will sooner or later limit scalability
- Large amounts of memory seldom used, but allocated for complete execution time
- Analogy to caches

✦ More sophisticated partitioning methods

- More messaging overhead
- Basic idea is based on the associativity of the addition
- Operate on blocks (or tiles) of the matrices that you own, then communicate to get new blocks






Parallel MMULT – 1D column partitioning

✦ Idea:

1. First, let i compute all it owns:
 - $C[:,i] = C[:,i] + A[:,i] * B[i,:]$
 $i = \text{my_rank}$
2. Then, process i sends $A[:,i]$ to process $(i+1) \% p$
 - p is number of processes
3. Repeat for $p-1$ times



```
C[:,m] = C[:,m] + A[:,m]*B[m,]
S = A[:,m]
```

```
For i = 1 to p-1
    send S to process (m+1)%p
    recv S from process (m-1)%p
    // S is now A[:,(m-i)%p]
```

```
C[:,m] = C[:,m] + S*B[(m-i)%p,]
```



Parallel MMULT – 1D column partitioning

✦ Computation costs per step

- p processes $elements * madd * (vector\ elements / blocks)$
- Computation costs per step: $f = N^2 * 2 * (N / p) = 2N^3/p$

✦ Timing model for communication

- Communication follows a linear model: $t(s) = t_{lat} + s * t_{BW}$
- s : size (number of elements), t_{lat} : constant overhead in time, t_{BW} : time to transfer one element

✦ Communication costs per step

- $s = N^2/p$ (size of A divided by number of processes)

✦ Total costs:

- $c = f + (p-1) * (f + t(s))$
 - Initial step compute, then $(p-1)$ times repeating
- $c = f * p + (p-1) * t_{lat} + (p-1) * N^2/p * t_{BW}$



Parallel MMULT – 1D column partitioning

✦ Total costs $c_{\text{parallel}} = 2N^3/p + (p-1)*t_{\text{lat}} + (1-1/p)*N^2 * t_{\text{BW}}$

- $t_{\text{new}} = c_{\text{parallel}}; t_{\text{old}} = f = 2N^3$

✦ Comments:

- Term 1: Perfect scalability ($O(1/p)$) if no communication costs, resp. no communication
- Term 2+3: Overhead due to parallelization

✦ Scaling the problem size

- $c_{\text{parallel}} = O(N^3/p)$
- $N \rightarrow \infty: SU \rightarrow ???$



Parallel MMULT – 1D column partitioning

✦ Total costs $c_{\text{parallel}} = \underbrace{2N^3/p}_{\text{Computation}} + \underbrace{(p-1)*t_{\text{lat}} + (1-1/p)*N^2 * t_{\text{BW}}}_{\text{Communication}}$

✦ Scalability

- Computation $O(N^3)$
- Communication $O(N^2)$

✦ As long as communication costs grow slower than computation costs, system is considered scalable

- One of the strictest definitions
- Rarely used because difficult to verify



✦ Parallel MMULT – 1D column partitioning

- Initialization: distribute $A[:,i]$ to process i , $i=1..P$, broadcast B
 - MPI_SCATTER
 - MPI_BCAST
- During each step: sending and receiving of blocks
 - MPI_SEND & MPI_RECV
 - MPI_SENDRECV
- Final result: gather $C[:,i]$ from process i , $i=1..P$
 - MPI_GATHER


✦ Collective operations can help to simplify a program, like library calls

- Every involved process has to make the function call



✦MPI_BCAST


- `MPI_Bcast (void *buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`
- As if process *root* executes *n* send operations to *n* destinations





✦MPI_SCATTER


- MPI_Scatter (void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)
- As if process *root* executes n send operations to n destinations
- See also MPI_Scatterv - various





✦MPI_GATHER

- MPI_Gather (void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)
- As if process *root* makes *n* receives from *n* sources
- See also MPI_Gatherv – various, MPI_Allgather (shown in gray)





Notes on collective operations

- ✦ Plenty of them, to cover all the various patterns
- ✦ Like libraries, others are responsible for optimization
 - E.g., for different topologies
- ✦ Communication cost: between $O(n)$ and $O(n^2)$
 - n =process count
- ✦ Dedicated hardware support for multicasts
 - Only found in advanced communication devices
 - BlueGene, Cray, EXTOLL
 - Typically not covered in specifications like Ethernet or Infiniband
- ✦ Multicast \neq Ethernet broadcast
 - Ethernet broadcasts are not reliable
- ✦ Currently collectives are always blocking (preventing overlap)
 - See MPI 3.0 for more advanced solutions



Parallel MMULT – 2D partitioning

✦ 1D partitioning

- Maps nicely to ring, also to mesh/tori/hypercube but does not utilize all network resources
- Still huge memory requirements (N^2/P)
- $P \sim N^2$ ⑨ scaling P with N is difficult
- $P \sim N$ would be better

✦ Assume $P=s^2$, s whole number

✦ Then each process computes a block of C :

- For all i, j of block(C); for $k = 0..(s-1)$
 - $C[i, j] = C[i, j] + \text{sum} (A [i, k] * B [k, j])$
- Rotate blocks of A & B circularly

✦ Owner computes, optimize locality



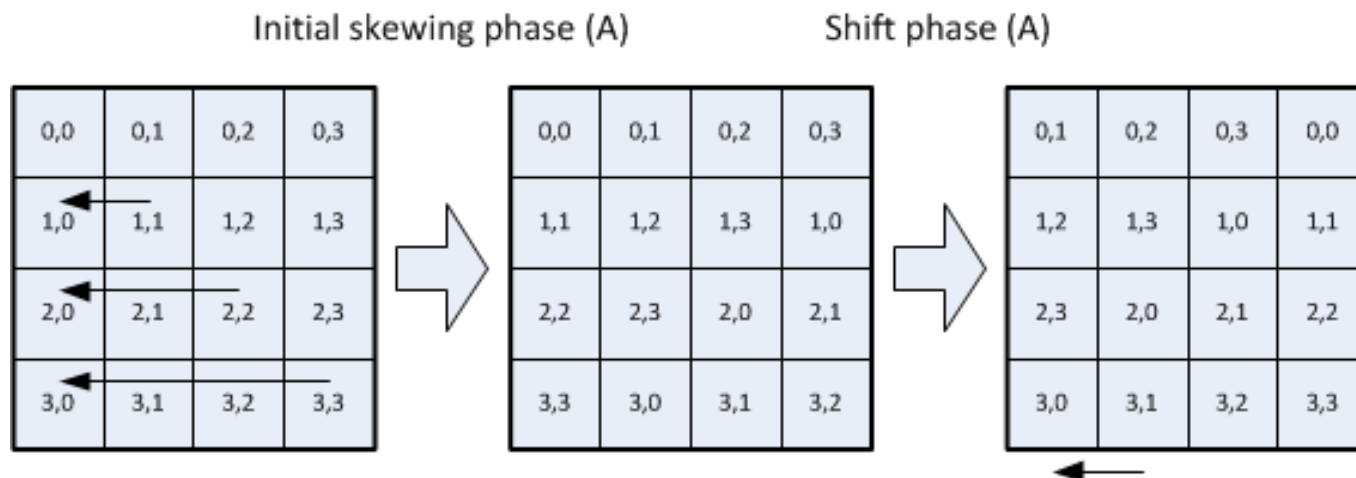
Parallel MMULT – 2D partitioning

✦ Three phases:

1. Initial skewing phase: left shift circular row i of A by i times to the left, up shift circular column j of B by j times
 - For all i : $A[i,j] = A[(i+1)\%s,j]$, B accordingly
2. Computation
3. Shift: {left,up} shift circular of each {row,column} of { A,B } by one

✦ In-place algorithm

✦ Communication costs now by factor \sqrt{p} smaller





✦ Broadcast-Multiply-Roll (BMR)

- Using a broadcast instead of the initial skewing process
- Three phases: broadcast, multiply, roll
- Algorithm is not in-place

✦ SUMMA: Scalable Universal Matrix Multiplication Algorithm

- Highly general algorithm
- B small: less memory requirement, lower efficiency; B large: high memory requirement, higher efficiency

✦ Many more

- Recursive layouts



✦ PCAM used for the MMULT example

- P: easy as domain decomposition is applied
- C: trivial (copies), difficult (1D/2D partitioning)
- A: trivial (copies), difficult (1D/2D partitioning)
- M: easy

✦ Data handling (initial distribution & later movement) is most difficult when writing MPI programs

✦ Sequential versions can help to verify correctness

✦ Collective MPI calls can make the program easier, both for the initial coder and later readers

- Topology agnostic