



Introduction to High Performance Computing

Lecture 02 – CUDA Programming


Holger Fröning
Institute of Computer Engineering
Ruprecht-Karls University of Heidelberg

With material from D. Kirk, W. Hwu („Programming Massively Parallel Processors“)




Introduction

GK110



XEON-E7






GPUs vs. CPUs

	Tesla K20	Xeon E7-4800 4P
Core count	13 SMs 64/832 (DP), 192/2,496 (SP)	10 Cores 2 FP-ALUs/core
Frequency	0.7GHz	2.4GHz
Peak Compute Performance	1,165 GFLOPS (DP) 3,494 GFLOPS (SP)	96 GFLOPS (DP)
Use model	throughput-oriented	latency-oriented
Latency treatment	toleration	minimization
Programming	1000s-10,000s of threads	10s of threads
Memory bandwidth	250 GBytes/sec	34 GByte/s (per P)
Memory capacity	<= 8 GB	up to 2TB
Die size	550mm ²	684 mm ²
Transistor count	7.1 billion	2.3 billion
Technology	28nm	32nm



Introduction

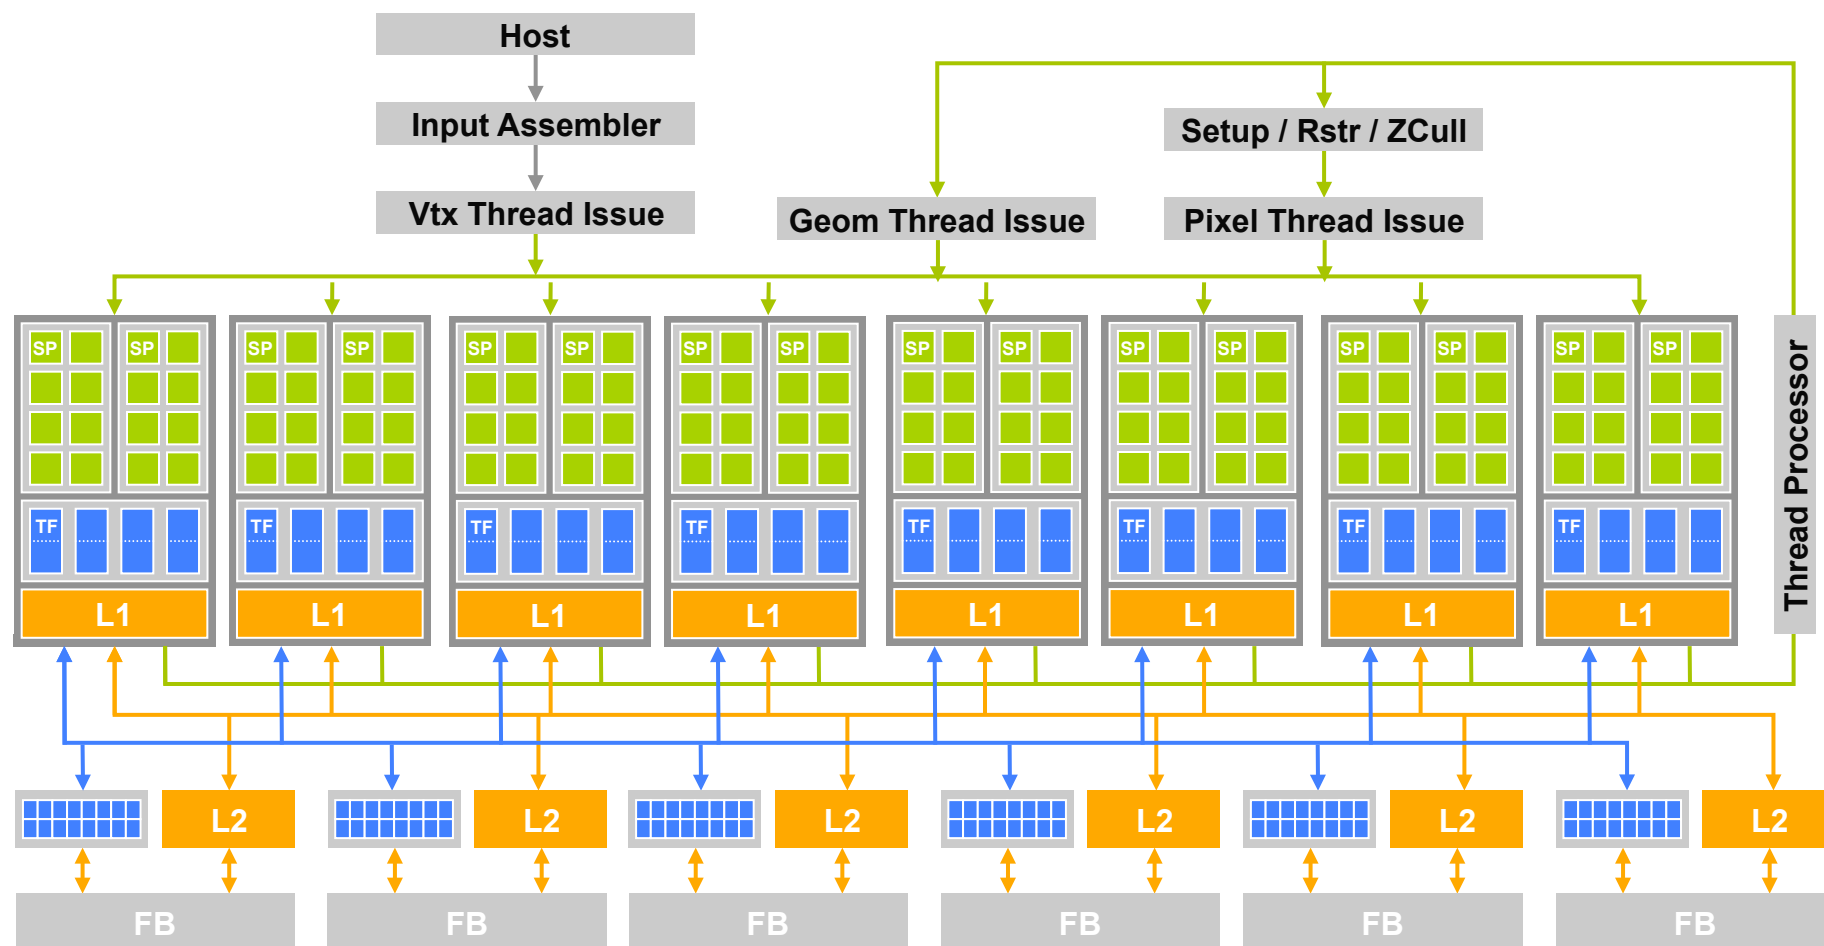
- **NVidia CUDA**
 - **Compute kernel** as C program
 - Explicit data- and thread-level parallelism
 - Computing, not graphics processing
 - Host communication
- **Memory hierarchy**
 - Registers at thread level
 - Shared memory at thread block level
 - GPU memory
 - Host memory
- **More HW details exposed**
 - Use of pointers
 - Load/store architecture
 - Barrier synchronization of thread blocks
- **Requires more detailed understanding of underlying HW**





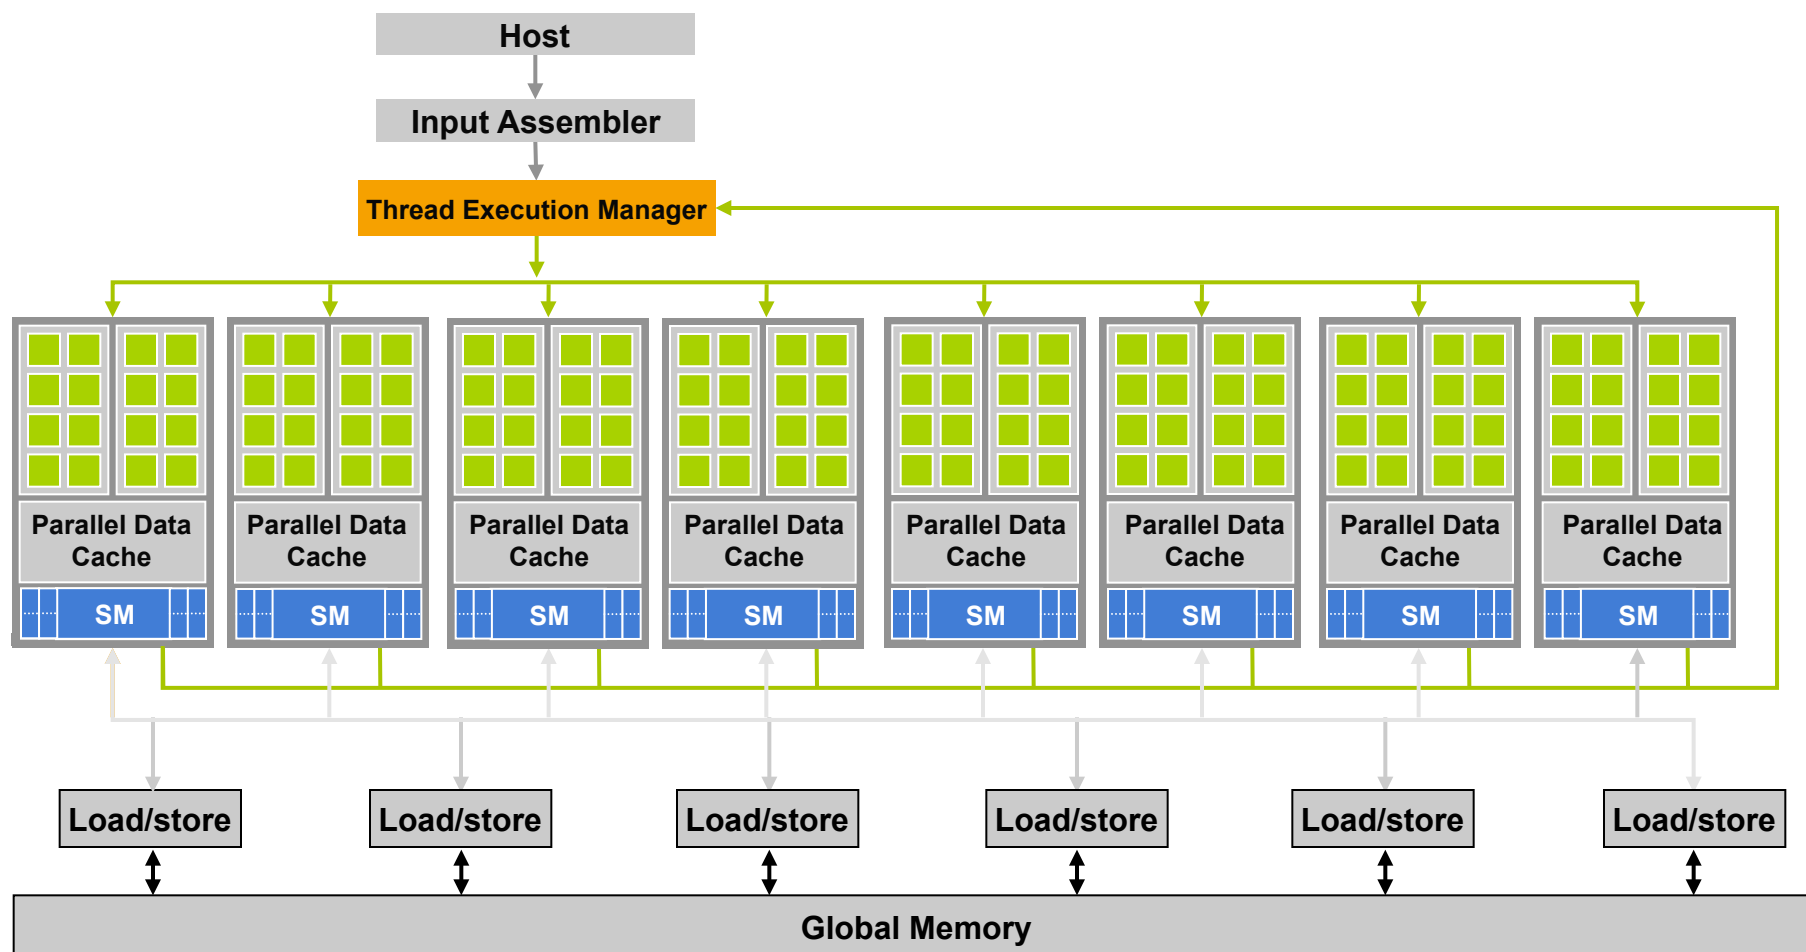
Introduction

■ G80 architecture for graphic processing





■ G80 architecture for general-purpose computing





CUDA Programming Model




Programming Model

■ C Extension with three main abstractions

1. Hierarchy of threads
2. Shared memory
3. Barrier synchronization

■ Exploiting parallelism

- Fine-grain DLP
- TLP




■ CUDA program consists of CPU & GPU part

- CPU part: part of the program with no or little parallelism
- GPU part: high parallel part, SPMD-style

■ Concurrent execution

- Non-blocking thread execution
- Explicit synchronization





Programming Model

```
__global__ void matAdd (float A[N][N], float B[N][N], float C[N][N])
{
    int i = threadIdx.x;
    int j = threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    // Kernel invocation
    dim3 dimBlock ( N, N );
    matAdd <<< 1, dimBlock >>> ( A, B, C );
}
```



- **Kernels:** n-fold execution by N threads
 - Definition: keyword `__global__`
- **Execution:** kernel `<<<NumBlocks, threadsPerBlock>>>` (args)
- **(Unique) ID:** `threadIdx`
 - Control flow for SPMD programs
 - Memory access orchestration





Programming Model

```
__global__ void matAdd (float A[N][N], float B[N][N], float C[N][N])
{
    int i = threadIdx.x;
    int j = threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    // Kernel invocation
    dim3 dimBlock ( N, N );
    matAdd <<< 1, dimBlock >>> ( A, B, C );
}
```

- `threadIdx` has up to 3 dimensions: `threadIdx.{x,y,z}`

➔ Each thread block has up to 3 dimensions

- Number of threads per block is limited

- 512 x 512 x 64 ➔ 1024 x 1024 x 64 (implementation dep.)

➔ Additional hierarchy level: grid = blocks of threads

- Unique ID `blockIdx`, up to 3 dimensions
- Blocks are executed independently and implementation-dependent
- Number of blocks limited (typ. 64k-1 per dimension)



up to 3D



Programming Model


```
__global__ void matAdd (float A[N][N], float B[N][N], float C[N][N])
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    if ( i < N && j < N )
        C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    // Kernel invocation
    dim3 dimBlock ( 16, 16 );
    dim3 dimGrid  ( ( N + dimBlock.x - 1 ) / dimBlock.x,
                    ( N + dimBlock.y - 1 ) / dimBlock.y );
    matAdd <<< dimGrid, dimBlock >>> ( A, B, C );
}
```

- Operator “/” rounds down, so add block size to round up!
- E.g. N=50:
- grid size = $(50+16-1)/16=4.0625 \Rightarrow 4$

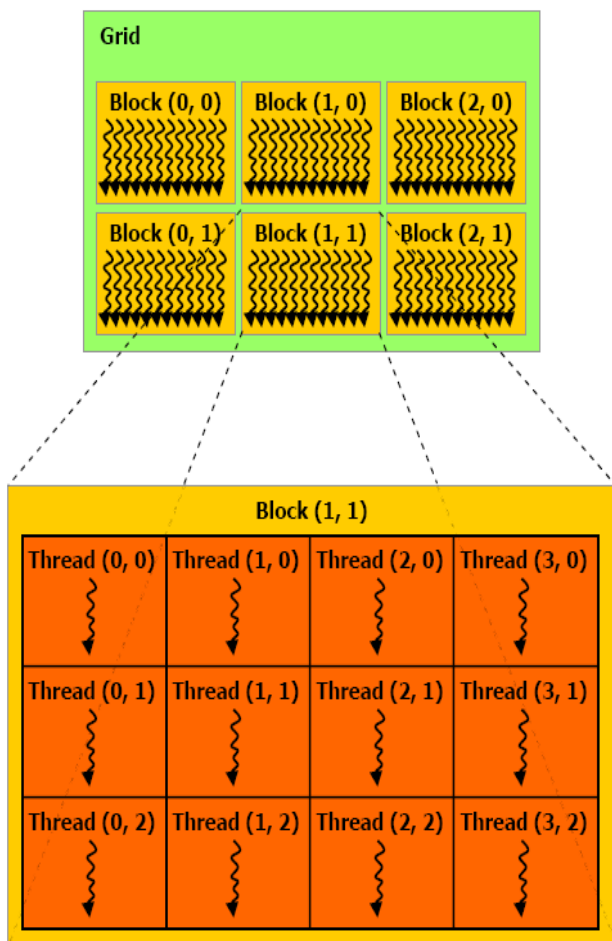
Block size

Fine grain PCAM
One thread computes
one element





Programming Model



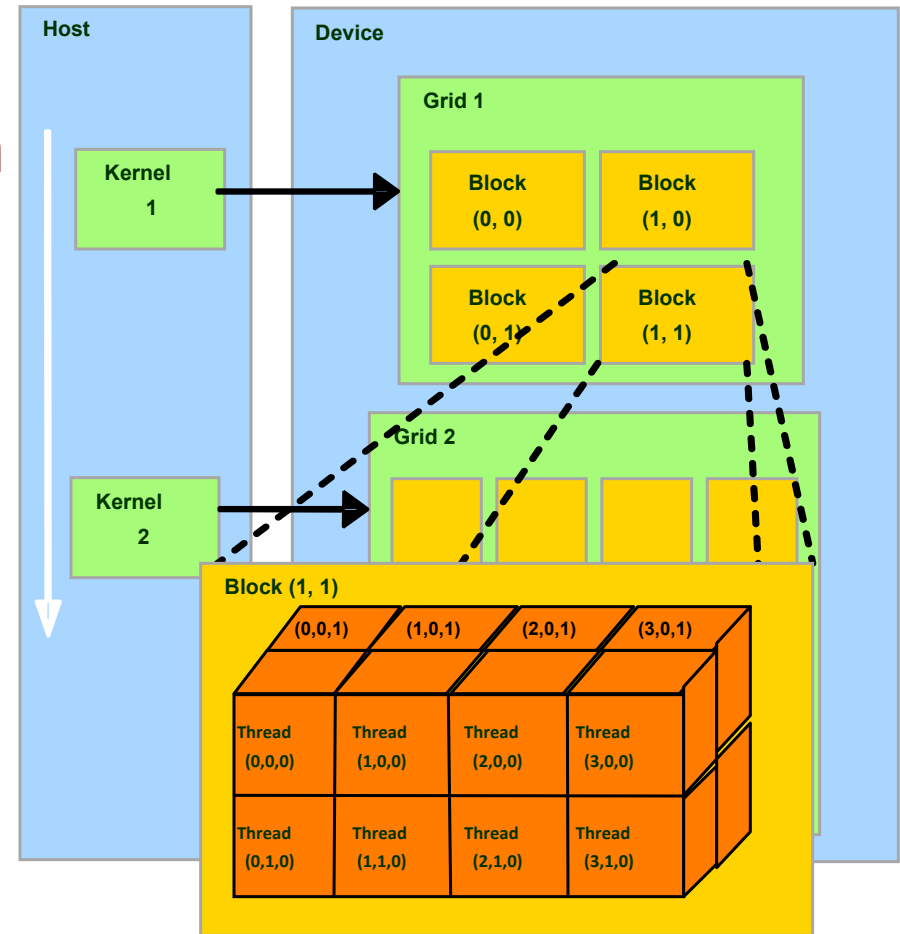
- **Thread hierarchy**
 - Grid of thread blocks
 - Blocks of equal size
- **Given problem size N, how to choose parameters threads per block respectively blocks per grid?**
- **Recommendations wrt block count**
 - >2x number of SMs
 - Optimal: 100 – 1000 (max. 64k-1)
- **Recommendations wrt threads/block**
 - Parallel slackness vs. number of registers per thread
 - $R / (B * \text{ceil}(T, 32))$

R = total registers, B = active blocks / SM, T = threads per block, ceil = round up to next multiple of 32
 R = 8k/SM .. 32k/SM (implementation dependent)



Programming Model

- Communication and synchronization only within one thread block
 - Shared memory
 - Atomic operations
 - Barrier synchronization
- Threads from different blocks cannot interact
 - Exception: global memory
 - Very weak coherence & consistency guarantees
- Iterative kernel invocations





Memory Hierarchy – Global Memory

■ Global memory

- Communication between host and device
- Accessible from all threads
 - Read/write
- High associated latency (no caching)
- Lifetime exceeds thread lifetime

■ Sensitive to coalescing

■ Allocation

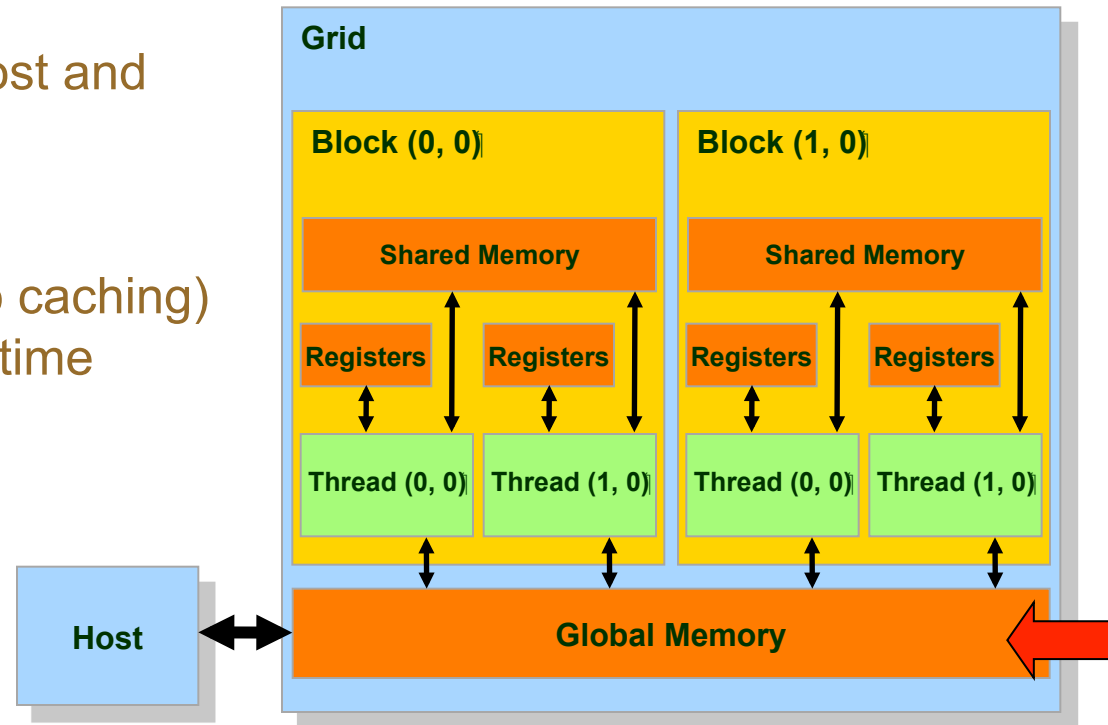
- `cudaMalloc (&dmem, size);`

■ De-allocation

- `cudaFree (dmem);`

■ Data transfer (blocking)

- `cudaMemcpy (*dst, *src, size, transfer_type);`
- `cudaMemcpyAsync (...)`





Memory Hierarchy – Global Memory

Variable scope annotation

```
void *dmem = cudaMalloc ( N*sizeof ( float ) ); // Allocate GPU memory

void *hmem = malloc ( N*sizeof ( float ) ); // Allocate CPU memory

// Transfer data from host to device
cudaMemcpy ( dmem, hmem, N*sizeof ( float ), cudaMemcpyHostToDevice );

// Do calculations
kernel1 <<< numBlocks, numThreadsPerBlock >>> ( dmem, N );
...
kernel2 <<< numBlocks, numThreadsPerBlock >>> ( dmem, N );

// Transfer data from device to host
cudaMemcpy ( hmem, dmem, N*sizeof ( float ), cudaMemcpyDeviceToHost );

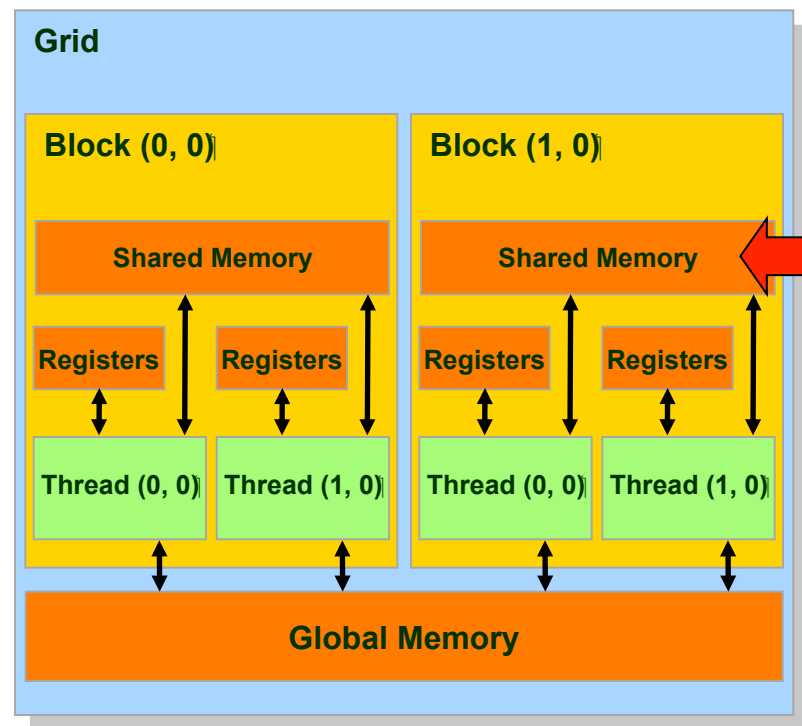
cudaFree ( dmem ); // Free device buffer
free ( hmem ); // Free host buffer
```

Only references
to device
memory



Memory Hierarchy – Shared Memory


- **Shared Memory**
 - On-chip memory
 - Lifetime: thread lifetime
- **Access costs in the best case equal register access**
 - Organized in n banks
 - Typ. 16-32 banks with 32bit width
 - Low-order interleaving
 - Parallel access if no conflict
 - Conflicts result in access serialization






Memory Hierarchy – Shared Memory

Shared Memory
Bank Access without Blocking



Shared Memory
Bank Access with Blocking





Memory Hierarchy - Declaration

	Location	Access from
<code>__device__ float var;</code>	global mem	device/host
<code>__constant__ float var;</code>	constant mem	device/host
<code>__shared__ float var; }</code>	shared mem	block

- **__device__** can be combined with one additional type
 - Otherwise placed in global memory, application lifetime
- **__constant__** implies **__device__**, see above
- **__shared__** implies **__device__**
 - Lifetime of a block, only available within this block
 - **__syncthreads** to wait for outstanding write operations
 - Unconstrained completion of read/write operations (exception: **volatile**)



■ Vector types

- char1, uchar1, char2, uchar2, char3, uchar3, char4, uchar4, short1, ushort1, short2, ushort2, short3, ushort3, short4, ushort4, int1, uint1, int2, uint2, int3, uint3, int4, uint4, long1, ulong1, long2, ulong2, long3, ulong3, long4, ulong4, float1, float2, float3, float4, double2
- Derived from basic types (int, float, ...)

■ Dimension type: dim3

- Based on uint3
- Unspecified components are initialized with 1



Function Declarations

	Executed on	Callable from
<code>__device__ float</code>	device	device
<code>DeviceFunc() ↵</code> <code>__global__ void KernelFunc() ↵</code>	device	host
<code>__host__ float HostFunc() ↵</code>	host	host

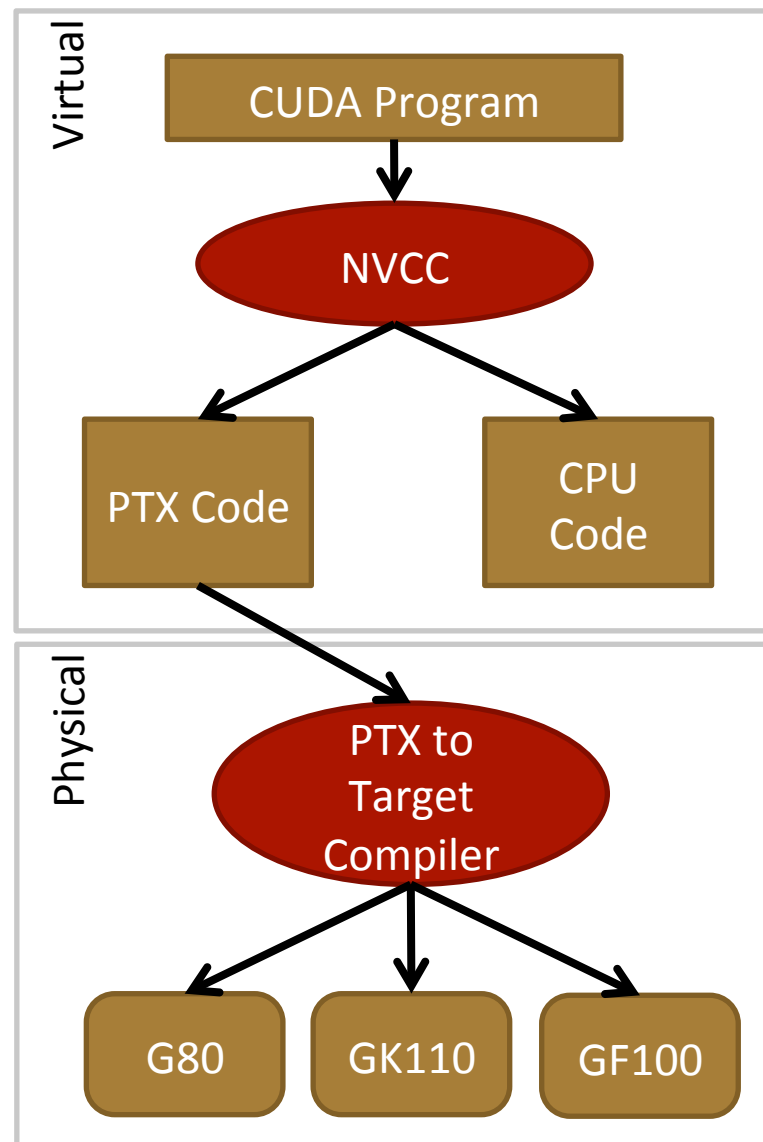
- `__global__` defines a kernel (return type: void)
 - Rückgabewert vom Typ void
- `__host__` is optional
- `__host__` and `__device__` can be combined
- No pointers to `__device__` functions
 - Exception: `__global__` functions
- For functions that are executed on the GPU:
 - No recursions, only static variable declarations, no variable parameter count



Just-in-time Compilation

- Device code only supports C-subset of C++
- Compile with **nvcc**
 - Compiler Driver
 - Calls other required tools
 - cudacc, g++, cl, ...
 - Debugging:
 - command line parameter -deviceemu
- **Output**
 - C code (host CPU Code)
 - Either PTX object code, or source code for run time interpretation
- **PTX (Parallel Thread Execution)**
 - Virtual Machine and ISA
 - Execution resources and state
- **Linking**
 - CUDA runtime library cudart
 - CUDA core library cuda

```
export LD_LIBRARY_PATH=/usr/local/cuda/lib
export PATH=$PATH:/usr/local/cuda/bin
nvcc -help
nvcc foo.cu -o foo
```





Brief Property Survey (deviceQuery)

Model	Revision number	Total global memory [bytes]	Multiprocessors	Cores	Total constant memory [bytes]	Shared memory per block [bytes]	Registers per block	Warp size	Threads per block	Max dimension of a block	Max. dimension of a grid	Max. memory pitch [bytes]	Clock rate [GHz]	Concurrent copy and execution
GeForce GTX 280	1.3	1G	30	240	64k	16k	16k	32	512	512 x 512 x 64	65535 x 65535 x 1	256k	1.3	Y 1
GeForce GTX 480	2.0	1.5G	15	480	64k	48k	32k	32	1k	1k x 1k x 64	65535 x 65535 x 65535	2G	1.4	Y 1
Tesla K20c	3.5	5G	13	2496	64k	48k	64k	32	1k	1k x 1k x 64	2G x 65535 x 65535	2G	0.7	Y 2



CUDA Example: saxpy



SAXPY Example

Scalar Alpha X Plus Y
 $y[i] = \text{alpha} * x[i] + y[i]$

- **SAXPY: Scalar Alpha X Plus Y**
- Simple test to compare the GPU and the CPU
 - Objective: runtime reduction
 - Max. Gridsize * threadsPerBlock elements
 - $65535 * 1k \rightarrow \sim 64M$ elements
 - Memory requirement = $32M \text{ elements} * 2 \text{ arrays} * 4 \text{ Byte/element} = 256MB$
- Source code contains kernels for the GPU and the CPU



SAXPY Example

```

////////////////////////////////////
// kernel function (CPU)
////////////////////////////////////
void saxpy_serial(int n, float alpha, float *x, float *y)
{
    int i;
    for (i=0; i<n; i++) {
        y[i] = alpha*x[i] + y[i];
    }
}

```

```

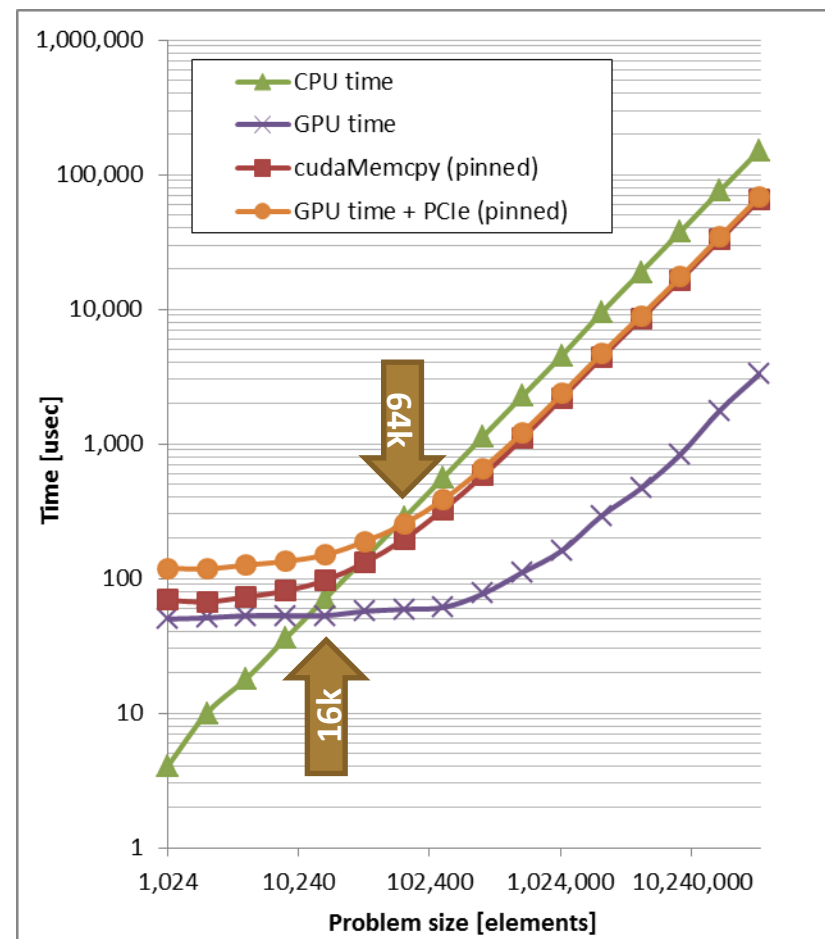
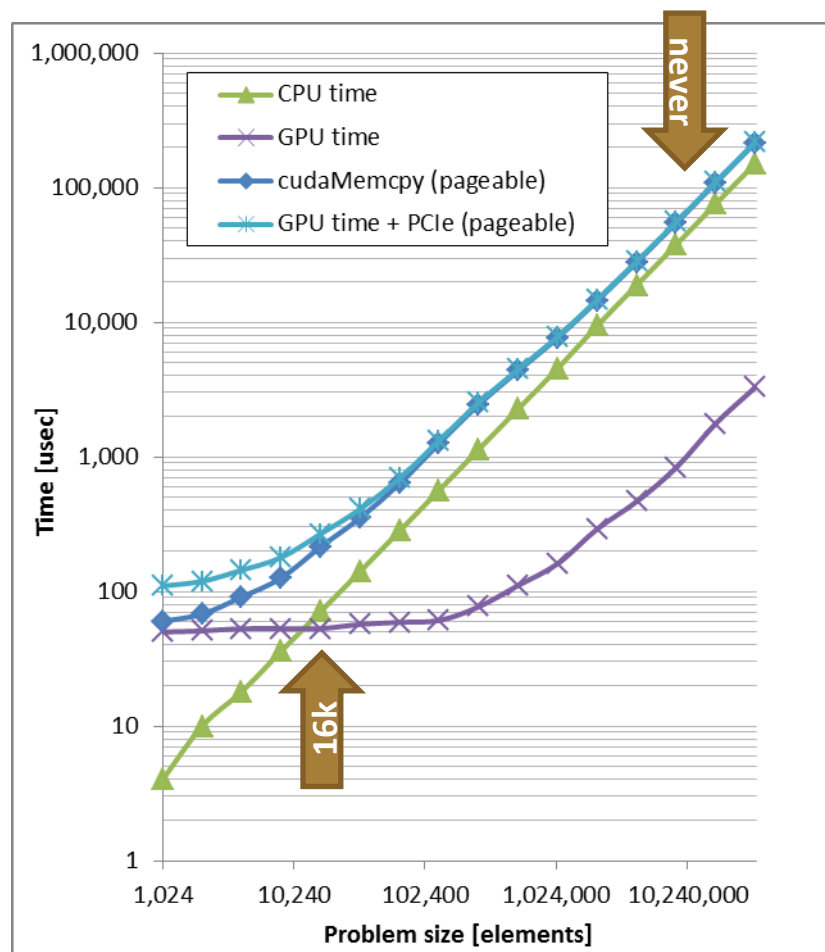
////////////////////////////////////
// kernel function (CUDA device)
////////////////////////////////////
__global__ void saxpy_parallel(int n, float alpha, float *x, float *y)
{
    // compute the global index in the vector from
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    // avoid writing past the allocated memory for the vector y.
    if (i<n) {
        y[i] = alpha*x[i] + y[i];
    }
}

```



SAXPY Example (2)



- Tesla K10, PCIe 2.0 x16 connection, Intel Xeon E5, single-threaded
- Break-even at 16k
- Huge advantage for GPU when no data movements
- Pageable memory vs. pinned memory



SAXPY Example (3)

- Replace malloc with cudaMallocHost
 - Significant reduce data movement costs
- Pinned memory is a scarce resource!

```

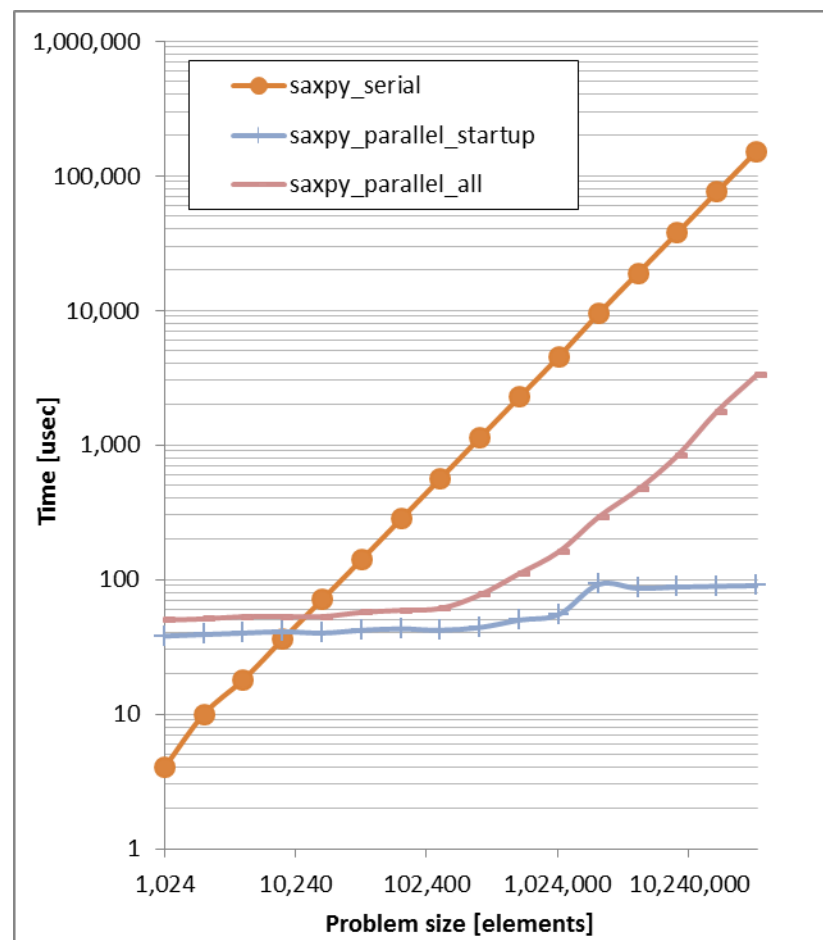
////////////////////////////////////
// (2) allocate memory on host (main CPU memory) and device,
//      h_ denotes data residing on the host, d_ on device
////////////////////////////////////
float *h_x;
float *h_y;
float *d_x;
float *d_y;
if (USE_PINNED_MEMORY) {
    cudaMallocHost ( (void**) &h_x, N*sizeof(float) );
    cudaMallocHost ( (void**) &h_y, N*sizeof(float) );
} else {
    h_x = (float*) malloc( N*sizeof(float) );
    h_y = (float*) malloc( N*sizeof(float) );
}
cudaMalloc((void**)&d_x, N*sizeof(float));
cudaMalloc((void**)&d_y, N*sizeof(float));
checkErrors("memory allocation");

```



SAXPY Example

- **Successful measures** 😊
 - Use of pinned memory
 - (Increase computational intensity)
- **Unsuccessful** 😞
 - Precompilation: 40usec kernel startup time is maintained, independent of pre-compiled code or not
- **SAXPY is a trivial example**
 - See dependency analysis





Summary – CUDA & GPU Computing

■ Introduction to CUDA

- Pretty unusual concept compared to CPU programming
- Pretty easy programming model
- Pretty challenging for upmost speed-ups

■ Direct control over hardware

- Plenty of opportunities for the (experienced) user
- Increases the burden

■ Main differences to CPU programming

- Sophisticated resource planning and usage
- Scratch pad: use shared memory as explicitly controlled cache
- Data movement over PCIe
- Limited memory

■ Key characteristics for good GPU applications

- High degree of DLP
- Low data reuse
- High computational intensity
- High bandwidth



- Runtime problems:

CUDA Error: the launch timed out and was terminated
→ Stop X11

CUDA Error: unspecified launch failure
→ Typically a segmentation fault

CUDA Error: invalid configuration argument
→ Too many threads per block, too many resources per thread (shared memory, register count)

- Compile problem:

mmult.cu(171): error: identifier "__eh_curr_region" is undefined

→ Non-static shared memory, use static allocation of shared memory