

Intro HPC: Blatt 4

17.11.1014

4.1 Reading

4.2 Matrix Multiply – parallel version using MPI

For the Matrix Multiply in parallel version we use the source code 'matrix_multiply.cc'. To run and compile use our file 'make+run.sh'.

In this code we first initialise the matrices A and B by the master process (rank==0). This process is then distributing the different lines of A to the different slave processes. The Matrix B is distributed to all processes completely. After the processes have calculated the result resulting of a line of A they are sending the result back to the master process. After this has got all the results it prints the solution. (This is not contribution to the calculation time).

The outcome of this calculation is the following Matrix C:

0	30	60	90	120
0	40	80	120	160
0	50	100	150	200
0	60	120	180	240
0	70	140	210	280

4.3 Matrix Multiply – scaling process count

We now execute our program with a different number of processes. For that we encounter the nodes creek[01-04]. We are calculating a Matrix of 1000x1000 entries. This we compare with our sequential implementation.

Plotting the Time that was needed for the calculations we see that it is decreasing very fast but is reaching limit very soon.

Plotting the speed-up we see an increase very similar to a linear speed up in the beginning but then it starts to deviate more and more.

The Efficiency is nearly 1 for about 4-8 processes but then drastically decreasing for higher numbers.

All this is a consequence of the process and the used data being at its best at about 4-8 processes. With higher values the message traffic gets too high, so that the parallelization is not very effective any more.

4.4 Matrix Multiply – scaling problem size

This can be changed if you have more data, so that the single process is having more work to do with nearly the same amount of message traffic.

