

Intro HPC: Blatt 7

24.11.1014

7.1 Reading: LogP: Towards a Realistic Model of Parallel Computation

The author is introducing a model, that describes the bottlenecks of a distributed memory multiprocessor in which processors communicate by point-to-point messages. Therefore it is based on a few parameters like latency (L), overhead (o), bandwidth (g) of communication and the number of processes (P) and the assumption of a finite capacity. Furthermore it is general for different types of communication protocols or applications.

The author is finding, that for some applications some parameters can get negligible, which makes a simplification of the model possible. He tested the model on different workloads like the FFT and the LU-Decomposition, to show how the use of the model can lead to efficient parallel applications in practice. In the paper he is also comparing this model to the widely used PRAM and the BSP model, which do not accurately reflect the performance characteristics of such systems, in his opinion.

This model is promising a couple of advantages to other often used models like including asynchronous algorithms. It is very general, so it might give a good overview of a system but not very detailed.

7.2 Reading: Roofline - An Insightful Visual Performance Model for Multicore Architectures

In this paper the author introduces a visual computational model for multicore Architectures called 'roofline' model. It relates the operational intensity (mean operations per byte of DRAM traffic) with an upper bound for performance of a kernel (Attainable GFlops/s). For that 'roofline' uses the minimum of Peak Floating-Point Performance (constant) and the Peak memory Bandwidth multiplied with the operational intensity (line with positive slope), so whether the problem is compute-bound or memory-bound.

In that way it tells the user what optimizations should be implemented and in what order, by pointing out the limiting factor of performance. Using micro benchmarks he is testing 4 different kernels to apply this model.

The Roofline model seems to be a very clear and easy to apply way for characterizing multicore architectures for to find a first starting point for optimizations.

7.3 n-Body Problem – Partitioning/Communication Design

to avoid collective calls the sum over all other particles in the force calculation starts at $i+1$ so firstly we make sure no pair is counted twice and secondly the firstly needed particle position is different for every process.

the masses are fixed so broadcast it once.